

Regression Analysis

Mohamed Shakir

December 8, 2023

Contents

1	Introduction	3
2	Simple Linear Regression	4
2.1	Estimating the Parameters with the Method of Least Squares	4
2.1.1	Estimation of β_0 and β_1	4
2.1.2	Properties of Fitted Regression Line	8
2.1.3	Estimation of σ^2	9
2.2	Hypothesis Testing on the Slope and Intercept	10
2.2.1	Using t -tests	10
2.2.2	Testing Significance	11
2.2.3	Analysis of Variance Tables (ANOVA)	11
2.3	Interval Estimation	12
2.3.1	Confidence Intervals on β_0 , β_1 , and σ^2	12
2.3.2	Interval Estimation of the mean Response	13
2.4	Prediction of New Observations	13
2.5	Coefficient of Determination	13
2.6	Correlation Coefficient	14
3	Multiple Linear Regression	15
3.1	Matrix Approach to Regression	15
3.1.1	Derivatives	16
3.2	Multiple Regression Models	16
3.2.1	Least Squares Estimation of Regression Coefficients	16
3.2.2	Properties of the Hat Matrix H	18
3.2.3	Properties of the Least-Squares Estimators	18
3.2.4	Estimation of σ^2	18
3.3	Estimation and Hypothesis Testing in Multiple Linear Regression	19
3.3.1	Testing for Significance of Regression	20
3.3.2	Tests on Individual Regression Coefficients	21
3.3.3	Extra Sum of Squares Principle	22
3.3.4	Testing the General Linear Hypothesis	23
3.4	Lack of Fit of the Regression Model	23
3.4.1	Test for Lack of Fit	24
3.5	Confidence Intervals in Multiple Regression	25
3.5.1	Confidence Intervals on Regression Coefficients	25
3.5.2	Confidence Intervals On the Mean Response	25
3.5.3	Simultaneous Confidence Intervals on Regression Coefficients	26
4	Model Adequacy	27
4.1	Residual Analysis	27
4.1.1	Checking for Normality	27
4.1.2	Checking Constant Variance	27
5	Transformations and Weighting to Correct Models	29
5.1	Variance Stabilizing Transformations	29
5.2	Transformations to Linearize the Model	29
5.2.1	Box-Cox Transformations	30
5.3	Generalized and Weighted Least Squares	31
5.3.1	Weighted Least Squares	31

5.4	Breusch-Pagen Test	32
6	Diagnostics for Leverage and Measures of Influence	33
6.1	Leverage	33
6.1.1	Properties of Leverage	33
6.2	Measures of Influence	34
6.2.1	Difference in Fit	34
6.2.2	Cook's Distance	34
6.2.3	Difference in Coefficients	34
7	Polynomial Regression Models	35
7.1	Polynomial Models in One Variable	35
7.2	Polynomial Models in Two Variables	36
7.3	Indicator Regression Models	36
8	Multicollinearity	37
8.1	Properties of Multicollinearity	37
8.2	Detecting Multicollinearity	37
8.2.1	Variance Inflation Factors	38
8.3	Ridge Regression	39
8.3.1	LASSO	40
9	Building the Regression Model	41
9.1	Criteria for Model Selection	41
9.1.1	Mallow's C_p	41
9.1.2	Akaike Information Criterion	42
9.1.3	Prediction Sum of Squares Criterion	43
9.2	Computational Techniques for Variable Selection	43
9.2.1	Forward Selection	43
9.2.2	Backward Elimination	43
9.2.3	Stepwise Regression	43
10	Logistic Regression and General Linear Models	44
10.1	Logistic Distribution	44
10.2	Estimating Parameters	45
10.2.1	Interpretation of the Parameters	45
10.2.2	Repeat Observations	45
10.3	Multiple Logistic Regression	46
10.3.1	Inference on Model Parameters	46
10.4	Testing	46
10.4.1	Test for Goodness of Fit	47
10.4.2	Hosmer-Lemeshow Test	47
10.4.3	Deviance of Goodness of Fit Test	47
10.5	Diagnostic Measures for Logistic Regression	48
10.6	Detection of Influential Observations	48
10.6.1	Influence on the Fitted Linear Predictor	48
10.7	Poisson Regression	48
10.8	Generalized Linear Models	49

Chapter 1

Introduction

The primary goal in regression is to develop a model that relates a set of explanatory variables X_1, \dots, X_p to a response variable Y , then test the model and use it for inference and prediction.

Given a set of n pairs of data Y_i and X_i , we attempt to fit a straight line to these points, using a simple regression model

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Where ϵ_i represents an unobserved random error term, β_0 is the intercept and β_1 is the slope of the line. β_0 and β_1 are parameters that need to be estimated from observed data. The model can also be expressed in terms of $(X_i - \bar{X})$.

$$Y_i = (\beta_0 + \beta_1 \bar{X}) + \beta_1 (X_i - \bar{X}) + \epsilon_i$$

Where \bar{X} is the sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

This proposed model is linear in the parameters β_0, β_1 , and would still be referred to as linear if we had X_i^2 instead of X_i . This model also makes the assumption that the random error terms ϵ_i are uncorrelated, have mean 0, and variance σ^2 . Under these assumptions, we have

$$E(Y_i) = \beta_0 + \beta_1 X_i$$

$$\text{Var}(Y_i) = \sigma^2$$

Chapter 2

Simple Linear Regression

The primary goal in regression is to develop a model that relates a set of explanatory variables X_1, \dots, X_p to a response variable Y , then test the model and use it for inference and prediction.

Given a set of n pairs of data Y_i and X_i , we attempt to fit a straight line to these points, using a simple regression model

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Where ϵ_i represents an unobserved random error term, β_0 is the intercept and β_1 is the slope of the line. β_0 and β_1 are parameters that need to be estimated from observed data. The model can also be expressed in terms of $(X_i - \bar{X})$.

$$Y_i = (\beta_0 + \beta_1 \bar{X}) + \beta_1 (X_i - \bar{X}) + \epsilon_i$$

Where \bar{X} is the sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

This proposed model is linear in the parameters β_0, β_1 , and would still be referred to as linear if we had X_i^2 instead of X_i . This model also makes the assumption that the random error terms ϵ_i are uncorrelated, have mean 0, and variance σ^2 . Under these assumptions, we have

$$E(Y_i) = \beta_0 + \beta_1 X_i$$

$$\text{Var}(Y_i) = \text{Var}(\beta_0 + \beta_1 X_i + \epsilon_i) = \sigma^2$$

Thus the mean of Y is a linear function of X however the variance of Y does not depend on a value of X .

The parameters β_0 and β_1 are called the regression coefficients. The slope β_1 is the change in the mean of the distribution of Y produced by a unit change in X . If the range of data on X includes $x = 0$, then the intercept β_0 is the mean of the distribution of the response Y when $x = 0$. If the range of x does not include zero, then β_0 has no practical interpretation.

2.1 Estimating the Parameters with the Method of Least Squares

The parameters β_0, β_1 are unknown and must be estimated from the data. Suppose we have n pairs of data $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$.

2.1.1 Estimation of β_0 and β_1

The method of least squares is the most popular approach to fitting a regression model. Let Q be the sum of the error terms squared

$$Q = \sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i)^2$$

Then we want to minimize Q with respect to the parameters β_1, β_2 ,

$$\frac{\partial Q}{\partial \beta_0} = -2 \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i) = 0$$

$$\frac{\partial Q}{\partial \beta_1} = -2 \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i) X_i = 0$$

We can rearrange these equations to get the following equations

$$\begin{aligned} & -2 \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i) = 0 \\ \implies & \sum_{i=1}^n Y_i - \sum_{i=1}^n \beta_0 - \sum_{i=1}^n \beta_1 X_i = 0 \\ \implies & \sum_{i=1}^n Y_i = n\beta_0 - \beta_1 \sum_{i=1}^n X_i \\ & -2 \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i) X_i = 0 \\ \implies & \sum_{i=1}^n Y_i X_i - \sum_{i=1}^n \beta_0 X_i - \sum_{i=1}^n \beta_1 X_i^2 = 0 \\ \implies & \sum_{i=1}^n Y_i X_i = \beta_0 \sum_{i=1}^n X_i - \beta_1 \sum_{i=1}^n X_i^2 \end{aligned}$$

These 2 equations are known as the normal equations and the solutions to them, call them b_0 , b_1 , are

$$b_0 = \bar{Y} - b_1 \bar{X}$$

$$b_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{\sum_{i=1}^n (X_i - \bar{X})Y_i}{\sum_{i=1}^n (X_i - \bar{X})^2} = \sum_{i=1}^n k_i Y_i$$

with

$$k_i = \frac{X_i - \bar{X}}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

We also sometimes use a more compact notation, by denoting the corrected sum of squares for X and the sum of cross products of X_i Y_i as

$$\begin{aligned} S_{xx} &= \sum_{i=1}^n x_i^2 - \frac{1}{n} \left(\sum_{i=1}^n x_i \right)^2 = \sum_{i=1}^n (x_i - \bar{x})^2 \\ S_{xy} &= \sum_{i=1}^n x_i y_i - \frac{1}{n} \left(\sum_{i=1}^n x_i \right) \left(\sum_{i=1}^n y_i \right) = \sum_{i=1}^n y_i (x_i - \bar{x}) \end{aligned}$$

So, we can write

$$b_1 = \frac{S_{xy}}{S_{xx}}$$

The observed difference between Y_i and the corresponding fitted value \hat{Y}_i is a residual. The i th residual is

$$e_i = y_i - \hat{y}_i = y_i - (b_0 + b_1 x_i)$$

Note that k_i has important properties, such as

$$\sum_{i=1}^n k_i = 0, \quad \sum_{i=1}^n k_i X_i = 1, \quad \sum_{i=1}^n k_i^2 = \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

$$\begin{aligned} \sum_{i=1}^n k_i &= \frac{\sum_{i=1}^n (X_i - \bar{X})}{\sum_{i=1}^n (X_i - \bar{X})^2} \\ &= \frac{n\bar{X} - n\bar{X}}{\sum_{i=1}^n (X_i - \bar{X})^2} = 0 \end{aligned}$$

$$\begin{aligned}
\sum_{i=1}^n k_i X_i &= \frac{\sum_{i=1}^n (X_i - \bar{X}) X_i}{\sum_{i=1}^n (X_i - \bar{X})^2} \\
&= \frac{\sum_{i=1}^n (X_i^2 - X_i \bar{X})}{\sum_{i=1}^n (X_i - \bar{X})^2} \\
&= \frac{\sum_{i=1}^n X_i^2 - \bar{X} \sum_{i=1}^n X_i}{\sum_{i=1}^n (X_i - \bar{X})^2} \\
&= \frac{\sum_{i=1}^n X_i^2 - n \bar{X}^2}{\sum_{i=1}^n (X_i^2 - 2X_i \bar{X} + \bar{X}^2)} \\
&= \frac{\sum_{i=1}^n X_i^2 - n \bar{X}^2}{\sum_{i=1}^n X_i^2 - 2n \bar{X}^2 + n \bar{X}^2} = 1
\end{aligned}$$

$$\begin{aligned}
\sum_{i=1}^n k_i^2 &= \sum_{i=1}^n \left(\frac{X_i - \bar{X}}{\sum_{i=1}^n (X_i - \bar{X})^2} \right)^2 \\
&= \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^4} \\
&= \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}
\end{aligned}$$

The equation for the fitted line is then

$$\hat{Y} = b_0 + b_1 X$$

Or alternatively using $X - \bar{X}$,

$$\hat{Y} = (b_0 + b_1 \bar{X}) + b_1 (X - \bar{X})$$

Theorem 2.1.1 (Gauss Markov Theorem). *The least square estimators b_0, b_1 are unbiased and have minimum variance among all unbiased linear estimators.*

Proof. Consider an unbiased linear estimator

$$\hat{\beta}_1 = \sum_{i=1}^n c_i Y_i$$

$\hat{\beta}_1$ must satisfy $E(\hat{\beta}_1) = \beta_1$.

$$\begin{aligned}
\beta_1 &= E(\hat{\beta}_1) \\
&= E\left(\sum_{i=1}^n c_i Y_i\right) \\
&= \sum_{i=1}^n c_i E(Y_i) \\
&= \sum_{i=1}^n c_i (\beta_0 + \beta_1 X_i) \\
&= \beta_0 \sum_{i=1}^n c_i + \beta_1 \sum_{i=1}^n c_i X_i
\end{aligned}$$

Therefore, $\sum_{i=1}^n c_i = 0$, and $\sum_{i=1}^n c_i X_i = 1$. We can also see that the variance is

$$\text{Var}(\hat{\beta}_1) = \sum_{i=1}^n c_i^2 \text{Var}(Y_i) = \sigma^2 \sum_{i=1}^n c_i^2$$

Now, set $c_i = k_i + d_i$ where k_i is as defined previously above and d_i are arbitrary constants. We want to show

that the variance is minimized, so

$$\begin{aligned}
\text{Var}(\hat{\beta}_1) &= \sum_{i=1}^n c_i^2 \text{Var}(Y_i) \\
&= \sigma^2 \sum_{i=1}^n c_i^2 \\
&= \sigma^2 \sum_{i=1}^n (k_i + d_i)^2 \\
&= \sigma^2 \left(\sum_{i=1}^n k_i^2 + 2 \sum_{i=1}^n k_i d_i + \sum_{i=1}^n d_i^2 \right)
\end{aligned}$$

Note that the variance of b_1 is

$$\text{Var}(b_1) = \text{Var} \left(\sum_{i=1}^n k_i Y_i \right) = \sigma^2 \sum_{i=1}^n k_i^2 = \sigma^2 \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

Now notice that there is a relationship between the variance of $\hat{\beta}_1$ and b_1 , namely that the variance of $\hat{\beta}_1$ is the same as b_1 plus an additional constants but these constants are indeed 0.

$$\begin{aligned}
\sum_{i=1}^n k_i d_i &= \sum_{i=1}^n k_i (c_i - k_i) \\
&= \sum_{i=1}^n k_i c_i - \sum_{i=1}^n k_i^2 \\
&= \sum_{i=1}^n c_i \frac{X_i - \bar{X}}{\sum_{i=1}^n (X_i - \bar{X})^2} - \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} \\
&= \frac{\sum_{i=1}^n c_i X_i - \sum_{i=1}^n c_i \bar{X}}{\sum_{i=1}^n (X_i - \bar{X})^2} - \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}
\end{aligned}$$

We know that $\sum_{i=1}^n c_i = 0$ and $\sum_{i=1}^n c_i X_i = 1$, so this becomes

$$\sum_{i=1}^n k_i d_i = \frac{1 - 0}{\sum_{i=1}^n (X_i - \bar{X})^2} - \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} = 0$$

Therefore,

$$\text{Var}(\hat{\beta}_1) = \sigma^2 \left(\sum_{i=1}^n k_i^2 + \sum_{i=1}^n d_i^2 \right)$$

Clearly the variance is minimized when $d_i = 0$ for all i , thus

$$\text{Var}(\hat{\beta}_1) = \sigma^2 \sum_{i=1}^n k_i^2 = \text{Var}(b_1)$$

Thus the least squares estimator b_1 has minimum variance along all unbiased estimators. □

We may write

$$\hat{Y} = b_0 + b_1 X$$

for the estimated or fitted line, and

$$e_i = Y_i - \hat{Y}_i$$

for the estimated i^{th} residual. The estimate for the variance σ^2 is then

$$\hat{\sigma}^2 = \frac{\sum_{i=1}^n e_i^2}{n - 2}$$

The estimate of the variance σ^2 is also known as the mean square error (MSE).

2.1.2 Properties of Fitted Regression Line

- (i) $\sum_{i=1}^n e_i = 0$. Recall that $\hat{Y} = b_0 + b_1 X = (b_0 + b_1 \bar{X}) + b_1(X - \bar{X})$, and

$$\bar{Y} = b_0 + b_1 \bar{X}$$

So $\hat{Y} = \bar{Y} + b_1(X - \bar{X})$, then

$$\begin{aligned} \sum_{i=1}^n e_i &= \sum_{i=1}^n (Y_i - \hat{Y}_i) \\ &= \sum_{i=1}^n Y_i - \sum_{i=1}^n \hat{Y}_i \\ &= \sum_{i=1}^n Y_i - \sum_{i=1}^n (\bar{Y} + b_1(X_i - \bar{X})) \\ &= n\bar{Y} - n\bar{Y} + b_1 \sum_{i=1}^n (X_i - \bar{X}) \\ &= n\bar{Y} - n\bar{Y} + b_1(n\bar{X} - n\bar{X}) = 0 \end{aligned}$$

- (ii) $\sum_{i=1}^n Y_i = \sum_{i=1}^n \hat{Y}_i$. This follows from the previous property since

$$\sum_{i=1}^n e_i = \sum_{i=1}^n Y_i - \sum_{i=1}^n \hat{Y}_i = 0 \implies \sum_{i=1}^n Y_i = \sum_{i=1}^n \hat{Y}_i$$

- (iii) $\sum_{i=1}^n X_i e_i = 0$. This can be shown from the definition

$$\begin{aligned} \sum_{i=1}^n X_i e_i &= \sum_{i=1}^n X_i (Y_i - \hat{Y}_i) \\ &= \sum_{i=1}^n X_i (Y_i - b_0 - b_1 X_i) \\ &= \sum_{i=1}^n X_i Y_i - b_0 \sum_{i=1}^n X_i - b_1 \sum_{i=1}^n X_i^2 \\ &= b_0 \sum_{i=1}^n X_i + b_1 \sum_{i=1}^n X_i - b_0 \sum_{i=1}^n X_i - b_1 \sum_{i=1}^n X_i \\ &= 0 \end{aligned}$$

This is significant because it tells us that the dot product between the vector of explanatory variables $\vec{X} = (X_1, \dots, X_n)^T$ is orthogonal to the vector of error terms $\vec{e} = (e_1, \dots, e_n)^T$, and from the previous property we get that

$$\vec{e} \cdot \mathbf{1}_n = \sum_{i=1}^n e_i = 0$$

Hence the vectors $\{\mathbf{1}_n, X - \bar{X}\mathbf{1}_n\}$ are linearly independent and form a basis of the estimation space.

(iv) By applying the Pythagorean Theorem to the previous property we get

$$\begin{aligned}
||Y||^2 &= ||\hat{Y}||^2 + ||Y - \hat{Y}||^2 \\
\sum_{i=1}^n Y_i^2 &= \sum_{i=1}^n \hat{Y}_i^2 + \sum_{i=1}^n e_i^2 \\
&= \sum_{i=1}^n \bar{Y}^2 + b_1^2 \sum_{i=1}^n (X_i - \bar{X})^2 + \sum_{i=1}^n e_i^2 \\
\Rightarrow \sum_{i=1}^n Y_i^2 - n\bar{Y}^2 &= b_1^2 \sum_{i=1}^n (X_i - \bar{X})^2 + \sum_{i=1}^n e_i^2 \\
\sum_{i=1}^n (Y_i - \bar{Y})^2 &= b_1^2 \sum_{i=1}^n (X_i - \bar{X})^2 + \sum_{i=1}^n (Y_i - \hat{Y}_i)^2
\end{aligned}$$

This shows us the the total sum of squares is equal to the regression sum of squares plus the error sum of squares.

(v) The point (\bar{X}, \bar{Y}) is on the fitted line.

(vi) The sum of residuals weighted by their corresponding fitted value is 0, that is

$$\sum_{i=1}^n y_i e_i = 0$$

(vii) Under the normality assumption, $e_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$. The method of maximum likelihood leads to the method of least squares.

$$L(\beta_0, \beta_1, \sigma^2) = \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n e_i^2 \right)$$

So maximizing $L(\beta_0, \beta_1, \sigma^2)$ is equivalent to minimizing $\sum e_i^2$.

(viii) The variance for \hat{y} is,

$$\text{Var}(\hat{Y}) = \sigma^2 \sum_{i=1}^n \left(\frac{1}{n} + k_i (X_i - \bar{X})^2 \right) = \sigma^2 \left(\frac{1}{n} + \frac{(X - \bar{X})}{S_{xx}} \right)$$

We can replace σ^2 with the mean square error as we will see in the next section to compute the sample variance of \hat{Y} .

2.1.3 Estimation of σ^2

We need to estimate σ^2 to test hypotheses and construct interval estimates pertinent to the regression model. Ideally we would like this estimate not to depend on the adequacy of the fitted model. This is only possible when there are several observations on y for at least one value of x , or when prior information concerning σ^2 is available. When this approach cannot be used, the estimate of σ^2 is obtained from the residual or error sum of squares.

$$SSE = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

We can substitute \hat{y}_i for $b_0 + b_1 x_i$ and simplify to get

$$SSE = \sum_{i=1}^n y_i^2 - n\bar{y}^2 - b_1 S_{xy}$$

Moreover, the correct sum of squares of the response variable is

$$SST = \sum_{i=1}^n y_i^2 - n\bar{y}^2 = \sum_{i=1}^n (y_i - \bar{y})^2$$

Thus,

$$SSE = SST - b_1 S_{xy}$$

The residual sum of squares has $n - 2$ degrees of freedom, because we reserve 2 degrees of freedom for the estimators b_0, b_1 . We will later show that the expected value for SSE is

$$E(SSE) = (n - 2)\sigma^2$$

So an unbiased estimator of σ^2 is

$$\hat{\sigma}^2 = \frac{SSE}{n - 2} = MSR$$

The quantity MSR is known as the **residual mean square**. The root of $\hat{\sigma}^2$ is known as the **standard error of regression**.

2.2 Hypothesis Testing on the Slope and Intercept

To perform hypotheses tests and construct confidence intervals, we require that we make the additional assumption that the model errors ϵ_i are normally distributed. Thus, the complete assumptions are that the errors are normally and independently distributed with mean 0 and variance σ^2 , written as $\{\epsilon_i\} \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$. We will discuss how these assumptions can be checked through residual analysis later.

Suppose that we have the model $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$, where $\{\epsilon_i\} \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$. Then

$$(a) \frac{b_1 - \beta_1}{se(b_1)} \sim t_{n-2} \text{ where } se^2(b_1) = \frac{MSE}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

$$(b) \frac{b_0 - \beta_0}{se(b_0)} \sim t_{n-2} \text{ where}$$

$$se^2(b_0) = MSE \left(\frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right)$$

(c) MSE is an unbiased estimate of σ^2 and is independent of b_0, b_1 . Furthermore

$$\frac{(n - 2)MSE}{\sigma^2} \sim \chi_{n-2}^2$$

Proof. Proof will be shown when we generalize this using matrices in later sections. □

2.2.1 Using t -tests

Suppose we want to test that the slope is equal to a constant, β , we have the hypotheses

$$H_0 : \beta_1 = \beta, \quad H_1 : \beta_1 \neq \beta$$

Since $\{\epsilon_i\} \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$, the observations y_i are normally distributed with $\beta_0 + \beta_1 x_i$ and variance σ^2 . Then, b_1 is a linear combination of the observations, so it is normally distributed with mean β_1 and variance σ^2/S_{xx} . Therefore, our test statistic becomes

$$Z_0 = \frac{b_1 - \beta}{\sqrt{\sigma^2/S_{xx}}}$$

If the null hypothesis is true, then $Z_0 \sim N(0, 1)$. If σ^2 is known then we would use Z_0 to test our hypotheses. However, σ^2 is typically unknown. We've seen that MSE is an unbiased estimator for σ^2 , and we've established that $(n - 2)MSE/\sigma^2 \sim \chi_{n-2}^2$.

$$t_0 = \frac{b_1 - \beta}{\sqrt{MSE/S_{xx}}}$$

If the null hypothesis is true, $t_0 \sim t_{n-2}$. We compare the observed value t_0 with the upper $\alpha/2$, of the t_{n-2} distribution. So we reject the null hypothesis

$$|t_0| > t_{\alpha/2, n-2}$$

We can also test with the p -value. From the equation for t_0 , the denominator is called the **estimated standard error** of the slope.

$$se(b_1) = \sqrt{\frac{MSE}{S_{xx}}}$$

So, we often write t_0 is

$$t_0 = \frac{b_1 - \beta}{se(\beta_1)}$$

We test the intercept in a similar manner,

$$H_0 : \beta_0 = \beta, \quad H_1 : \beta_0 \neq \beta$$

We use a similar test statistic,

$$t_0 = \frac{b_0 - \beta}{se(b_0)}$$

and we reject the null hypothesis when $|t_0| > t_{\alpha/2, n-2}$.

2.2.2 Testing Significance

A special case for hypotheses is

$$H_0 : \beta_1 = 0, \quad H_1 : \beta_1 \neq 0$$

These hypotheses relate to the **significance of regression**. Failing to reject the null hypothesis means there is no linear relationship between x and y , we would reject the null hypothesis when $|t_0| > t_{\alpha/2, n-2}$.

2.2.3 Analysis of Variance Tables (ANOVA)

Analysis of variance can be used to test significance of regression. The analysis of variance of variance is based on a partitioning of the total variability of the response variable y , given by

$$y_i - \bar{y} = (\hat{y}_i - \bar{y}) + (y_i - \hat{y}_i)$$

Then, taking the sum of the square of both sides

$$\sum_{i=1}^n (y_i - \bar{y})^2 = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + \sum_{i=1}^n (y_i - \hat{y}_i)^2 + 2 \sum_{i=1}^n (\hat{y}_i - \bar{y})(y_i - \hat{y}_i)$$

Notice that

$$\begin{aligned} 2 \sum_{i=1}^n (\hat{y}_i - \bar{y})(y_i - \hat{y}_i) &= 2 \sum_{i=1}^n \hat{y}_i (y_i - \hat{y}_i) - 2\bar{y} \sum_{i=1}^n (y_i - \hat{y}_i) \\ &= 2 \sum_{i=1}^n \hat{y}_i e_i - 2\bar{y} \sum_{i=1}^n e_i = 0 \end{aligned}$$

Therefore,

$$\sum_{i=1}^n (y_i - \bar{y})^2 = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

The left side is the corrected sum of squares of the observations, which we denote by SST or $SSTO$. Notice that $y_i - \hat{y}_i = e_i$, so that term is the sum of residuals squared SSE . We call $\sum_{i=1}^n (\hat{y}_i - \bar{y})^2$ the **regression sum of squares**. So we have

$$SST = SSR + SSE$$

The regression sum of squares can also be computed by

$$SSR = b_1^2 S_{xx}$$

The **degrees of freedom** for each sum of squares is as follows.

- The total sum of squares SST has $df_T = n - 1$ since we lose a degree of freedom for the constraint

$$\sum_{i=1}^n (y_i - \bar{y}) = 0$$

- The regression sum of squares SSR has $df_R = p - 1$ where p is the number of variables (including y).

- The residual sum of squares SSE has $df_E = n - 2$ degrees of freedom since 2 constraints are placed on $e_i = y_i - \hat{y}_i$ with the estimation for β_0 and β_1 .

We create a table to summarize our results from statistical analysis.

Source	SS	DF	MS=SS/df	E(MS)
Regression	$SSR = b_1^2 \sum (X_i - \bar{X})^2$	$p - 1$	MSR	$\sigma^2 + \beta_1^2 \sum (X_i - \bar{X})^2$
Error	$SSE = \sum (Y_i - \hat{Y}_i)^2$	$n - p$	MSE	σ^2
Total	$SSTO = \sum (Y_i - \bar{Y})^2$	$n - 1$		

Each of the sums of squares is a quadratic form where the rank of the corresponding matrix is the degrees of freedom indicated. Chochran's theorem applies and we conclude that the quadratic forms are independent and have chi-squared distributions. Note that

$$\frac{SSR}{\sigma^2} = \frac{b_1^2 \sum (X_i - \bar{X})^2}{\sigma^2} \sim \chi_{p-1}^2$$

$$\frac{SSE}{\sigma^2} = \frac{\sum (Y_i - \hat{Y}_i)^2}{\sigma^2} \sim \chi_{n-p}^2$$

Then, the ratio between 2 chi-squared distributions divided by their degrees of freedom has a F-distribution with their respective degrees of freedom.

$$F = \frac{SSR/\sigma^2(p-1)}{SSE/\sigma^2(n-p)} = \frac{SSR/(p-1)}{SSE/(n-p)} = \frac{MSR}{MSE} \sim F_{p-1, n-p}$$

The degrees of freedom are determined by the amount of data required to calculate each expression. To summarize, the ANOVA table indicates how one can test the null hypothesis

$$H_0 : \beta_1 = 0$$

$$H_1 : \beta_1 \neq 0$$

The null Hypothesis is that the slope of the line is equal to 0. Under the null hypothesis, the expected mean square for regression and the expected mean square error are separate independent estimates of the variance σ^2 .

2.3 Interval Estimation

2.3.1 Confidence Intervals on β_0 , β_1 , and σ^2 .

The width of the confidence intervals are a measure of the quality of the regression line. If the error is normally and independently distributed by our assumption, then $(b_1 - \beta_1)/se(b_1)$ and $(b_0 - \beta_0)/se(b_0)$ follow a t distribution with $n - 2$ degrees of freedom. So, a $100(1 - \alpha)$ percent. The confidence interval for the slope β_1 is

$$b_1 - t_{\alpha/2, n-2} se(b_1) \leq b_1 \leq b_1 + t_{\alpha/2, n-2} se(b_1)$$

and for the intercept β_0 ,

$$b_0 - t_{\alpha/2, n-2} se(b_0) \leq b_0 \leq b_0 + t_{\alpha/2, n-2} se(b_0)$$

The interpretation for these intervals is, if we were to take repeated samples of the same size at the same x levels and construct 95% CIs on the slope for each sample, then 95% of those intervals will contain the true value of β_1 .

As we've seen earlier, the sampling distribution of $(n - 2)MSE/\sigma^2$ follows a chi-square distribution with $n - 2$ degrees of freedom, so

$$P \left\{ \chi_{1-\alpha/2, n-2}^2 \leq \frac{(n-2)MSE}{\sigma^2} \leq \chi_{\alpha/2, n-2}^2 \right\}$$

Thus the $100(1 - \alpha)$ percent CI on σ^2 is

$$\frac{(n-2)MSE}{\chi_{\alpha/2, n-2}^2} \leq \sigma^2 \leq \frac{(n-2)MSE}{\chi_{1-\alpha/2, n-2}^2}$$

2.3.2 Interval Estimation of the mean Response

Another important part of the regression model is estimating the mean response $E(y)$ for a particular regressor variable x . Assuming that x_0 is any value of the regressor variable within the range of the original data on x that we used to create the model. Then, an unbiased estimator for $E(y|x_0)$ can be found from the fitting model

$$\widehat{E(y|x_0)} = \hat{\mu}_{y|x_0} = b_0 + b_1x_0$$

Note that $\hat{\mu}_{y|x_0}$ follows a normal distribution since it is a linear combination of the observations y_i . The variance is

$$\text{Var}(\hat{\mu}_{y|x_0}) = \text{Var}(b_0 + b_1x_0) = \text{Var}(\bar{y} - b_1(x_0 - \bar{x})) = \frac{\sigma^2}{n} + \frac{\sigma^2(x_0 - \bar{x})^2}{S_{xx}}$$

The sampling distribution for

$$\frac{\hat{\mu}_{y|x_0} - E(y|x_0)}{\sqrt{MSE(1/n + (x_0 - \bar{x})^2/S_{xx})}}$$

is a t distribution with $n - 2$ degrees of freedom. Then the CI is given as

$$\left[\hat{\mu}_{y|x_0} \pm t_{\alpha/2, n-2} \sqrt{MSE \left(\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{S_{xx}} \right)} \right]$$

2.4 Prediction of New Observations

When we want to predict a new value for our regressor variable, say $x = x_0$, we obtain a point estimate for the response y , given as

$$\hat{y}_0 = \hat{\beta}_0 + \hat{\beta}_1x_0$$

Now we want to obtain an interval estimate for our new observation y_0 and conduct hypothesis tests. Note that the confidence interval for the mean response at $x = x_0$ is **not** the same. We define the new random error variable

$$\psi = y_0 - \hat{y}_0$$

which is normally distributed with mean zero and variance

$$\text{Var}(\psi) = \text{Var}(y_0 - \hat{y}_0) = \sigma^2 \left(1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{S_{xx}} \right)$$

Then, we use the standard error for ψ for our prediction interval, giving us the $100(1 - \alpha)$ percent prediction interval

$$\left[\hat{y}_0 \pm t_{\alpha/2, n-2} \sqrt{MSE \left(1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{S_{xx}} \right)} \right]$$

We can also conduct hypothesis testing, suppose our hypotheses are $H_0 : y_0 = y_{00}$, $H_1 : y_0 \neq y_{00}$, then we use the test statistic

$$\frac{\hat{y}_0 - y_{00}}{\sqrt{MSE \left(1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{S_{xx}} \right)}} = \frac{\hat{y}_0 - y_{00}}{se(\psi)} \sim t_{n-2}$$

We reject the null hypothesis when $|t_0| > t_{\alpha/2, n-2}$.

2.5 Coefficient of Determination

The quantity

$$R^2 = \frac{SSR}{SST} = 1 - \frac{SSE}{SST}$$

is called the **coefficient of determination**. R^2 is also called the proportion of variation explained by the regressor x since SST is a measure of variability in y without considering the effect of x , and SSE is the variability in y after considering x . Since $0 \leq SSE \leq SST$, then $0 \leq R^2 \leq 1$. An R^2 value close to 1 means **most of the variability of y is explained by x** .

2.6 Correlation Coefficient

The pearson correlation coefficient, denoted by ρ , related to b_1 is given as

$$\rho = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)}\sqrt{\text{Var}(Y)}}$$

This measures the linear correlation between 2 variables. When applied to a sample,

$$r = b_1 \left(\frac{S_{xx}}{SST} \right)^{\frac{1}{2}} = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}} = \frac{S_{xy}}{(S_{xx}SST)^{1/2}}$$

Note that $-1 \leq r \leq 1$. To test hypotheses on ρ , we have 2 cases. The hypotheses for testing if the correlation is 0 is as follows

$$H_0 : \rho = 0, \quad H_1 : \rho \neq 0$$

When testing the null hypothesis $\rho = 0$, we use a t statistic given as

$$t = \frac{r\sqrt{n-2}}{\sqrt{1-r^2}} \sim t_{n-2}$$

When testing

$$H_0 : \rho = \rho_0, \quad H_1 : \rho \neq \rho_0$$

We use a Z statistic,

$$Z = \text{arctanh } r = \frac{1}{2} \ln \left(\frac{1+r}{1-r} \right) \sim N \left(\mu_z, \frac{1}{n-3} \right)$$

where

$$\mu_z = \frac{1}{2} \ln \left(\frac{1+\rho}{1-\rho} \right)$$

Now we can standardize our statistic to obtain a standard normal test statistic

$$Z_0 = (\text{arctanh}(r) - \text{arctanh}(\rho_0))\sqrt{n-3}$$

We can obtain our confidence interval with

$$\tanh \left(\text{arctanh}(r) - \frac{Z_{\alpha/2}}{\sqrt{n-3}} \right) \leq \rho \leq \tanh \left(\text{arctanh}(r) + \frac{Z_{\alpha/2}}{\sqrt{n-3}} \right)$$

where $\tanh(u) = (e^u - e^{-u})/(e^u + e^{-u})$.

Chapter 3

Multiple Linear Regression

We call a regression model with more than one regressor variable a **multiple regression model**.

3.1 Matrix Approach to Regression

We will first cover simple linear regression in matrix form.

Let $Y = [Y_1, \dots, Y_n]^T$ be a column data vector, and we'll define the expected value as

$$E(Y) = \begin{bmatrix} E(Y_1) \\ \vdots \\ E(Y_n) \end{bmatrix}$$

Proposition 3.1.1. *If $Z = AY + B$ for a matrix of constants A , and B , then*

$$E(Z) = AE(Y) + B$$

Proof. Simply from the definition of expectations on vectors,

$$E(Z_i) = E\left(\left[\sum_j a_{ij} Y_j\right] + b_i\right) = \sum_j a_{ij} E(Y_j) + b_i$$

So

$$E(Z) = AE(Y) + B$$

□

Definition 3.1.1. *The covariance of a vector of data*

$$Y = \begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix}$$

is

$$\text{Cov}(Y) = E([Y - E(Y)][Y - E(Y)]^T) = \Sigma$$

Proposition 3.1.2. $\text{Cov}(AY) = A\Sigma A^T$.

Definition 3.1.2. *A random vector Y has a multivariate normal distribution if its density is given by*

$$f(y_1, \dots, y_n) = \frac{|\Sigma|^{-1/2}}{\exp\left(-\frac{1}{2}(Y - \mu)^T \Sigma^{-1}(Y - \mu)\right)}$$

where

$$Y^T = (y_1, \dots, y_n), \quad \mu^T = (\mu_1, \dots, \mu_n)$$

we denote this by

$$Y \sim N_n(\mu, \Sigma)$$

Theorem 3.1.1. Let $Y \sim N_n(\mu, \Sigma)$. Let A be an arbitrary $p \times n$ matrix of constants. Then

$$Z = AY + B \sim N_p(A\mu + B, A\Sigma A^T)$$

This theorem implies that any linear combination of normal variates has a normal distribution. This theorem won't be proved here.

3.1.1 Derivatives

- $z = a'y \rightarrow \frac{\partial z}{\partial y} = a$
- $z = y'y \rightarrow \frac{\partial z}{\partial y} = 2y$
- $z = a' Ay \rightarrow \frac{\partial z}{\partial y} = A'a$
- $z = y' Ay \rightarrow \frac{\partial z}{\partial y} = A'y + Ay$
- If A is symmetric, then $z = y' Ay \rightarrow \frac{\partial z}{\partial y} = 2A'y$

3.2 Multiple Regression Models

Suppose we have 2 regressor variables, a multiple regression model that may describe a relationship with our data is

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$$

The parameter β_1 indicates the expected change in response per unit change in x_1 when x_2 is held constant. Similarly β_2 measures the change in y per unit change in x_2 when x_1 is held constant.

3.2.1 Least Squares Estimation of Regression Coefficients

The method of **least squares** can be used to estimate the regression coefficients. Suppose $n > k$ observations are available, and let y_i denote the i th observed response x_{ij} denote the i th observation or level of regressor. We can write the sample regression model as

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik} + \epsilon_i = \beta_0 + \sum_{j=1}^k \beta_j x_{ij} + \epsilon_i$$

for $i = 1, 2, \dots, n$. Then the least squares function is

$$S(\beta_0, \beta_1, \dots, \beta_k) = \sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^k \beta_j x_{ij} \right)^2$$

Similar to the simple linear regression approach with Q , we want to minimize S with respect to $\beta_0, \beta_1, \dots, \beta_k$. So the least squares estimators must satisfy

$$\left. \frac{\partial S}{\partial \beta_0} \right|_{\hat{\beta}_0, \dots, \hat{\beta}_k} = -2 \sum_{i=1}^n \left(y_i - \hat{\beta}_0 - \sum_{j=1}^k \hat{\beta}_j x_{ij} \right) = 0$$

and similarly for the rest of the estimators for β_j for $j = 1, \dots, k$.

$$\left. \frac{\partial S}{\partial \beta_j} \right|_{\hat{\beta}_0, \dots, \hat{\beta}_k} = -2 \sum_{i=1}^n \left(y_i - \hat{\beta}_0 - \sum_{j=1}^k \hat{\beta}_j x_{ij} \right) x_{ij} = 0$$

When simplifying these equations, we get the least squares **normal equations**, which we will put in terms of matrices later on.

$$\begin{aligned} n\hat{\beta}_0 + \hat{\beta}_1 \sum_{i=1}^n x_{i1} + \hat{\beta}_2 \sum_{i=1}^n x_{i2} + \cdots + \hat{\beta}_k \sum_{i=1}^n x_{ik} &= \sum_{i=1}^n y_i \\ \hat{\beta}_0 \sum_{i=1}^n x_{i1} + \hat{\beta}_1 \sum_{i=1}^n x_{i1}^2 + \hat{\beta}_2 \sum_{i=1}^n x_{i1}x_{i2} + \cdots + \hat{\beta}_k \sum_{i=1}^n x_{ik}x_{i1} &= \sum_{i=1}^n x_{i1}y_i \end{aligned}$$

$$\begin{aligned} & \vdots \\ & \hat{\beta}_0 \sum_{i=1}^n x_{ik} + \hat{\beta}_1 \sum_{i=1}^n x_{i1}x_{ik} + \hat{\beta}_2 \sum_{i=1}^n x_{ik}x_{i2} + \cdots + \hat{\beta}_k \sum_{i=1}^n x_{ik}^2 = \sum_{i=1}^n x_{ik}y_i \end{aligned}$$

Note that there are $p = k + 1$ normal equations for each of the unknown regression coefficients. The solutions give us the **least squares estimators** $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_k$.

As mentioned earlier, it is easier to work with matrix notation, so we will express the model using matrices

$$Y = X\beta + \epsilon$$

where

$$Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \quad X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1k} \\ 1 & x_{21} & x_{22} & \cdots & x_{2k} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{nk} \end{bmatrix}$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_k \end{bmatrix}, \quad \epsilon = \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}$$

In general, Y is an $n \times 1$ vector of observations, X is an $n \times p$ matrix of the levels of the regressor variables, β is a $p \times 1$ vector of regression coefficients, and ϵ is an $n \times 1$ vector of random variables. Note that from here on out we will denote the transpose of a matrix A^T as A' . We want to find the vector of least squares estimators $\hat{\beta}$ that minimizes

$$S(\beta) = \sum_{i=1}^n \epsilon_i^2 = \epsilon' \epsilon = (Y - X\beta)'(Y - X\beta)$$

This can be expanded as

$$S(\beta) = Y'Y - \beta'X'Y - Y'X\beta + \beta'X'X\beta = Y'Y - 2\beta'X'Y + \beta'X'X\beta$$

This is derived from using the fact that $\beta'X'Y$ is a 1×1 matrix, so it is a scalar and its transpose is the same. Now our equation must satisfy

$$\left. \frac{\partial S}{\partial \beta} \right|_{\hat{\beta}} = -2X'Y + 2X'X\hat{\beta} = 0$$

This gives us

$$X'X\hat{\beta} = X'Y$$

The equations here are our **least-squares** normal equations. They are the same as the previous equations we found earlier not in matrix form.

We can solve the normal equations by multiplying both sides of our equation by $X'X$ to give us

$$\hat{\beta} = (X'X)^{-1}X'Y$$

provided that the inverse exists.

The fitted regression model corresponding to the levels of regressor variables $x' = [1, x_1, x_2, \dots, x_k]$ is then

$$\hat{y} = x'\hat{\beta} = \hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j x_j$$

So the vector of fitted values \hat{y}_i corresponding to the observed values y_i is

$$\hat{Y} = X\hat{\beta} = X(X'X)^{-1}X'Y = HY$$

The $n \times n$ matrix $H = X(X'X)^{-1}X'$ is known as the **hat matrix**. It maps the vector of observed values to a vector of fitted values (in other words, it "puts the hat" on Y).

3.2.2 Properties of the Hat Matrix H

The hat matrix has some useful properties, notably

- (a) H is a projection matrix, so it is idempotent and symmetric

$$HH = H$$

$$H' = H$$

- (b) The matrix H is orthogonal to the matrix $I - H$, so

$$(I - H)H = H - HH = 0$$

Moreover, $(I - H)$ is idempotent and a project matrix as well.

- (c) The vector of residuals, which we will denote as \vec{e} , is given as

$$\vec{e} = Y - \hat{Y} = Y - HY = (I - H)Y$$

- (d) Properties (b) and (c) imply that the observation vector Y is projected onto a space spanned by the columns of H , and the residuals are in an orthogonal space (similar to the case for the simple linear regression model).

$$Y = HY + (I - H)Y$$

- (e) Similar to our simple linear model, we can apply the pythagorean theorem (from the fact that the matrices are orthogonal) to obtain

$$\|Y\|^2 = \|HY\|^2 + \|(I - H)Y\|^2$$

We will see later why expressing

$$\vec{e} = Y - X\hat{\beta} = Y - HY = (I - H)Y$$

is useful.

3.2.3 Properties of the Least-Squares Estimators

We can first show that $\hat{\beta}$ is an unbiased estimators,

$$\begin{aligned} E(\hat{\beta}) &= E[(X'X)^{-1}X'Y] \\ &= E[(X'X)^{-1}X'(X\beta + \epsilon)] \\ &= E[(X'X)^{-1}X'X\beta + X'\epsilon] \\ &= E[\beta + \epsilon] \\ &= E(\beta) + E(\epsilon) \\ &= \beta \end{aligned}$$

This is using the fact that $E(\epsilon) = 0$. So if our model assumptions hold, $\hat{\beta}$ is an unbiased estimator for β . The variance of $\hat{\beta}$ is expressed by the covariance matrix

$$\text{Cov}(\hat{\beta}) = E\left\{[\hat{\beta} - E(\hat{\beta})][\hat{\beta} - E(\hat{\beta})]^T\right\}$$

3.2.4 Estimation of σ^2

We may develop an estimator for σ^2 using the residual sum of squares,

$$SSE = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n e_i^2 = e'e$$

Then, we can substitute $e = Y - X\hat{\beta}$,

$$\begin{aligned} SSE &= (Y - X\hat{\beta})'(Y - X\hat{\beta}) \\ &= Y'Y - \hat{\beta}'X'Y - Y'X\hat{\beta} + \hat{\beta}'X'X\hat{\beta} \\ &= Y'Y - 2\hat{\beta}'X'Y + \hat{\beta}'X'X\hat{\beta} \end{aligned}$$

Then since $X'X\hat{\beta} = X'y$, we get

$$SSE = Y'Y - \hat{\beta}'X'Y$$

It can be shown that SSE has $n - p$ degrees of freedom since p parameters are estimated in the regression model. This gives us the **residual mean square** or **mean square error**

$$MSE = \frac{SSE}{n - p}$$

It can be shown that MSE is again an unbiased estimator for σ^2 , therefore

$$\hat{\sigma}^2 = MSE$$

3.3 Estimation and Hypothesis Testing in Multiple Linear Regression

The true relationship between y and our regressors x_1, \dots, x_k is unknown, and our multiple linear regression model is used to approximate this. Models sometimes are more complex in structure than the model we've discussed,

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \dots + \beta_kx_k + \epsilon$$

but we can still use a multiple linear regression model. For example, consider a cubic polynomial model

$$y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \epsilon$$

If we set $x_1 = x$, $x_2 = x^2$, and $x_3 = x^3$, we can rewrite the equation as

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \epsilon$$

Models that include **interaction effects** may also be analyzed by multiple linear regression methods. For example if we have the model

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_{12}x_1x_2 + \epsilon$$

we can set $x_3 = x_1x_2$, and $\beta_3 = \beta_{12}$, then we get

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \epsilon$$

Finally, consider a **second-order model with interaction**

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_{11}x_1^2 + \beta_{22}x_2^2 + \beta_{12}x_1x_2 + \epsilon$$

We can set $x_3 = x_1^2$, $x_4 = x_2^2$, $x_5 = x_1x_2$, $\beta_3 = \beta_{11}$, $\beta_4 = \beta_{22}$, and $\beta_5 = \beta_{12}$, then we can get our linear model

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \beta_4x_4 + \beta_5x_5 + \epsilon$$

Our predictor variables may also be qualitative, for example

$$X = \begin{cases} 0 & \text{if subject is male} \\ 1 & \text{if subject is female} \end{cases}$$

We could also have a transformed response variable

$$\ln Y_i = \beta_0 + \beta_1X_{i1} + \beta_2X_{i2} + \epsilon_i$$

Now we want to be able to ask questions about our models adequacy and which regressors are important. To do this, we can conduct various hypothesis tests which also require that our random errors are independent and identically distributed from a normal distribution with mean $E(\epsilon_i) = 0$, and variance $\text{Var}(\epsilon_i) = \sigma^2$.

3.3.1 Testing for Significance of Regression

The test for **significance of regression** is to determine if there is a linear relationship between the response Y and any of the regressors x_1, \dots, x_k . The appropriate hypotheses are

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0, \quad H_1 : \beta_j \neq 0$$

for *at least one* value of j . Rejection of the null hypothesis implies *at least one* of the regressors contribute significantly to the model. The test is a generalization of the analysis of variance used in simple regression. We define the sum of squares the same way,

$$SST = SSR + SSE$$

If the null hypothesis is true, then $SSR/\sigma^2 \sim \chi_k^2$, where k is the number of regressors in the model ($p - 1$). Similarly, $SSE/\sigma^2 \sim \chi_{n-k-1}^2$ if the null hypothesis holds and it can be shown that SSR/σ^2 is independent of SSE/σ^2 . We then use an F statistic

$$F_0 = \frac{SSR/k}{SSE/(n-k-1)} = \frac{MSR}{MSE} \sim F_{k, n-k-1}$$

which follows an F distribution. Then, we reject the null hypothesis $H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0$ if

$$F_0 > F_{\alpha, k, n-k-1}$$

We can construct an analysis of variance table to summarize this procedure again.

Source of Variation	Sum of Squares	DF	Mean Square	F -statistic
Regression	SSR	k	MSR	MSR/MSE
Residuals	SSE	$n - k - 1$	MSE	
Total	SST	$n - 1$		

The **total sum of squares** is

$$SST = \sum_{i=1}^n Y_i^2 - \frac{1}{n} \left(\sum_{i=1}^n Y_i \right)^2 = Y'Y - \frac{1}{n} \left(\sum_{i=1}^n Y_i \right)^2$$

The **regression sum of squares** is

$$SSR = \hat{\beta}' X' Y - \frac{1}{n} \left(\sum_{i=1}^n Y_i \right)^2$$

The **residual sum of squares** is

$$SSE = Y'Y - \hat{\beta}' X' Y = Y'(I - H)Y$$

We can show that $Y'J_nY = (\sum_{i=1}^n Y_i)^2$ where J_n is the matrix with 1 in every entry

$$\begin{aligned}
Y'J_nY &= \begin{bmatrix} Y_1 & \cdots & Y_n \end{bmatrix} \begin{bmatrix} 1 & \cdots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \cdots & 1 \end{bmatrix} \begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix} \\
&= \begin{bmatrix} \sum Y_i & \cdots & \sum Y_i \end{bmatrix} \begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix} \\
&= Y_1 \sum_{i=1}^n Y_i + Y_2 \sum_{i=1}^n Y_i + \cdots + Y_n \sum_{i=1}^n Y_i \\
&= \sum_{i=1}^n Y_i (Y_1 + Y_2 + \cdots + Y_n) \\
&= \left(\sum_{i=1}^n Y_i \right) \left(\sum_{i=1}^n Y_i \right) \\
&= \left(\sum_{i=1}^n Y_i \right)^2
\end{aligned}$$

Thus we can write

$$\begin{aligned}
SST &= \sum_{i=1}^n Y_i^2 - \frac{1}{n} \left(\sum_{i=1}^n Y_i \right)^2 \\
&= Y'Y - \frac{1}{n} Y'J_nY \\
&= Y' \left(Y - \frac{1}{n} J_nY \right) \\
&= Y' \left(I - \frac{1}{n} J_n \right) Y
\end{aligned}$$

It can be shown easily that we can write the regression sum of squares in terms of this matrix J_n

$$SSR = Y' \left(H - \frac{1}{n} J_n \right) Y$$

3.3.2 Tests on Individual Regression Coefficients

Once we have determined that at least one of the regressors is significant, we can investigate which one (or more) is significant. Adding a variable to the model always causes the regression sum of squares to increase and the residual sum of squares to decrease. So we must decide whether the increase in SSR is sufficient to warrant using an additional regressor. The added regressor also increases the variance for the fitted value \hat{y} , so we must be careful to add only regressors that are significant. Adding an unimportant regressor can increase the mean square error, which can decrease the usefulness of the model.

To test an individual regression coefficient, say β_j , we use the test

$$H_0 : \beta_j = 0, \quad H_1 : \beta_j \neq 0$$

If the null hypothesis is not rejected, then we can remove the corresponding regressor x_j . The **test statistic** for this hypothesis is

$$t_0 = \frac{\hat{\beta}_j}{\sqrt{\hat{\sigma}^2 C_{jj}}} = \frac{\hat{\beta}_j}{se(\hat{\beta}_j)}$$

Where C_{jj} is the diagonal element of the matrix $(X'X)^{-1}$ corresponding to β_j . We reject the null hypothesis if

$$|t_0| > t_{\alpha/2, n-k-1}$$

3.3.3 Extra Sum of Squares Principle

We may also directly determine the contribution to the regression sum of squares of a regressor, say x_j , by using the **extra sum of squares method**. This procedure can also be used to investigate the contribution of a subset of the regressor variables.

Consider the regression model with k regressors,

$$Y = X\beta + \epsilon$$

where Y is $n \times 1$, X is $n \times p$, β is $p \times 1$, and ϵ is $n \times 1$ with $p = k + 1$. We want to determine if some subset of $r < k$ regressors contribute significantly to our model. We will partition the regression coefficients into 2 vectors, β_1 is a $(p - r) \times 1$ vector, and β_2 is the $r \times 1$ vector of coefficients we are trying to test. We want to test the following hypotheses

$$H_0 : \beta_2 = 0, \quad H_1 : \beta_2 \neq 0$$

So our model will be rewritten as

$$Y = X\beta + \epsilon = X_1\beta_1 + X_2\beta_2 + \epsilon$$

where X_1 is a $n \times (p - r)$ matrix that are the columns of X associated with β_1 , and X_2 is an $n \times r$ matrix with the columns of X associated with β_2 . This is our **full model**.

For the full model, we have established that

$$\hat{\beta} = (X'X)^{-1}X'Y$$

and the regression sum of squares is

$$SSR(\beta) = \hat{\beta}'X'Y$$

which has $k = p - 1$ degrees of freedom. We also have the residual mean square

$$MSE = \frac{Y'Y - \hat{\beta}'X'Y}{n - p} = \frac{SSE}{n - p}$$

To find the contribution of the regressors in β_2 , we fit the model assuming the null hypothesis $H_0 : \beta_2 = 0$ is true. This gives us the **reduced model**

$$Y = X_1\beta_1 + \epsilon$$

The least squares estimator for β_1 in the reduced model is

$$\hat{\beta}_1 = (X_1'X_1)^{-1}X_1'Y$$

The regression sum of squares is

$$SSR(\beta_1) = \hat{\beta}_1'X_1'Y$$

which has $k - r = p - 1 - r$ degrees of freedom. The regression sum of squares due to β_2 given that β_1 is already in the model is

$$SSR(\beta_2|\beta_1) = SSR(\beta) - SSR(\beta_1)$$

with $(p - 1) - (p - 1 - r) = r$ degrees of freedom. This sum of squares is called the **extra sum of squares due to β_2** because it measures the increase in the regression sum of squares that results from adding $x_{k-r+1}, x_{k-r+2}, \dots, x_k$ to the model with x_1, x_2, \dots, x_{k-r} . Now we have the extra sum of squares due to β_2 is independent of mean square error, so the null hypothesis $\beta_2 = 0$ can be tested with the following statistic,

$$F_0 = \frac{SSR(\beta_2|\beta_1)/r}{MSE}$$

If $F_0 > F_{\alpha, r, n-p}$, then we reject the null hypothesis and conclude that at least one of the parameters in β_2 is not zero, and at least one of the regressors in X_2 contributes significantly to the regression model. This is sometimes called a **partial F test** since it measures the contribution of regressors in X_2 given that X_1 is in the model.

Alternatively, we can compute this using the residual sum of squares, we'll denote the full model residual sum of squares with $SSE(F)$ and its corresponding degrees of freedom as $df_F = n - p$, similarly for the reduced model we will use $SSE(R)$ and $df_R = n - p + r$. Note that $df_R - df_F = n - p + r - n + p = r$, and our test statistic becomes

$$F_0 = \frac{\frac{SSE(R) - SSE(F)}{df_R - df_F}}{\frac{SSE(F)}{df_F}} = \frac{(SSE(R) - SSE(F))/r}{MSE}$$

If we want to compute the regression squares for multiple variables, say

$$SSR(\beta_2|\beta_1, \beta_0)$$

We compute it in the following way

$$SSR(\beta_2|\beta_1, \beta_0) = SSR(\beta_2, \beta_1, \beta_0) - SSR(\beta_1, \beta_0) = SSR(\beta_1, \beta_2|\beta_0) - SSR(\beta_1|\beta_0)$$

We can partition the regression sum of squares into marginal single degree of freedom components. For example, consider the model

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \epsilon$$

We can use the following identity,

$$SST = SSR(\beta_1, \beta_2, \beta_3|\beta_0) + SSE$$

We can then decompose the three degree of freedom regression sum of squares as

$$SSR(\beta_1, \beta_2, \beta_3|\beta_0) = SSR(\beta_1, \beta_0) + SSR(\beta_2|\beta_1, \beta_0) + SSR(\beta_3|\beta_2, \beta_1, \beta_0)$$

3.3.4 Testing the General Linear Hypothesis

Suppose that the null hypothesis we want to test is $H_0 : T\beta = 0$ where T is an $r \times p$ matrix such that only r of the p equations in $T\beta = 0$ are independent. That is, the rows are the independent equations and will yield a $r \times 1$ vector of coefficients we'd like to test. Recall the β is a $p \times 1$ vector. The **full model** stays the same, with $Y = X\beta + \epsilon$ and $\hat{\beta} = (X'X)^{-1}X'Y$, and the residual sum of squares is

$$SSE(FM) = Y'Y - \hat{\beta}'X'Y$$

which has $n - p$ degrees of freedom. To obtain the **reduced model**, the r independent equations in $T\beta = 0$ are used to solve for r of the regression coefficients in the full model in terms of the $p - r$ regression coefficient. This gives us the reduced model

$$Y = Z\gamma + \epsilon$$

where Z is an $n \times (p - r)$ matrix, and γ is a $(p - r) \times 1$ vector of the unknown regression coefficients. The estimator for γ is

$$\hat{\gamma} = (Z'Z)^{-1}Z'Y$$

and the residual sum of squares is

$$SSE(RM) = Y'Y - \hat{\gamma}'Z'Y$$

which has $n - p + r$ degrees of freedom. The reduced model has less parameters than the full model so $SSE(RM) \geq SSE(FM)$. Now to test the hypothesis $H_0 : T\beta = 0$, we use the difference in residual sum of squares, denoted by SSH ,

$$SSH = SSE(RM) - SSE(FM)$$

which has $n - p + r - (n - p) = r$ degrees of freedom. This is known as the sum of squares due to the hypothesis. The test statistic here is then

$$F_0 = \frac{SSH/r}{SSE(FM)/(n - p)} \sim F_{r, n-p}$$

We reject H_0 if $F_0 > F_{\alpha, r, n-p}$.

3.4 Lack of Fit of the Regression Model

Sometimes, we want to test whether or not our linear model is justified at all. This is different from testing if the slope is 0.

3.4.1 Test for Lack of Fit

Suppose we have n_i observations on the response at the i th level of the regressor x_i for $i = 1, 2, \dots, m$. Let y_{ij} denote the j th ($j = 1, 2, \dots, n_i$) observation on the corresponding response variable. In total, there are

$$n = \sum_{i=1}^m n_i$$

observations. We partition the residual sum of squares

$$SSE = SS_{PE} + SS_{LOF}$$

Where SS_{PE} is the sum of squares due to **pure error** and SS_{LOF} is the sum of squares due to **lack of fit**. Note that we can partition SSE using the following,

$$y_{ij} - \hat{y}_i = (y_{ij} - \bar{y}_i) + (\bar{y}_i - \hat{y}_i)$$

\bar{y}_i is the average from the n_i observations at x_i , then we can square and sum both sides of the equation to give us

$$\sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \hat{y}_i)^2 = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_i)^2 + \sum_{i=1}^m \sum_{j=1}^{n_i} (\bar{y}_i - \hat{y}_i)^2$$

This gives us that the pure error sum of squares is

$$SS_{PE} = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_i)^2$$

Since there are $n_i - 1$ degrees of freedom for pure error at each level x_i that gives us the total degrees of freedom for SS_{PE} as

$$\sum_{i=1}^m (n_i - 1) = n - m$$

The sum of squares due to lack of fit is

$$SS_{LOF} = \sum_{i=1}^m n_i (\bar{y}_i - \hat{y}_i)^2$$

This is a weighted sum of squared deviations between the mean response \bar{y}_i at each level x_i with its corresponding fitted value \hat{y}_i . If the fitted values are close to the average responses, then there is a strong indication the regression function is linear. There are $m - 2$ degrees of freedom for SS_{LOF} since there are m levels and 2 degrees are reserved for estimating the 2 parameters β_0, β_1 to compute \hat{y}_i . This gives us our test statistic

$$F^* = \frac{SS_{LOF}/(m - 2)}{SS_{PE}/(n - m)} = \frac{MS_{LOF}}{MS_{PE}}$$

The expected value for MS_{PE} is σ^2 , and for MS_{LOF} it is

$$E(MS_{LOF}) = \sigma^2 + \frac{\sum_{i=1}^m n_i [E(y_i) - \beta_0 - \beta_1 x_i]^2}{m - 2}$$

If the regression function is linear, then $E(y_i) = \beta_0 + \beta_1 x_i$ so $E(MS_{LOF})$ is σ^2 . So, we construct our hypotheses as

$$H_0 : E(y_i) = \beta_0 + \beta_1 x_i, \quad H_1 : E(y_i) \neq \beta_0 + \beta_1 x_i$$

and we reject the null hypothesis if $F^* > F_{\alpha, m-2, n-m}$. We can summarize these in another analysis of variance table,

Source	Sum of Squares	DF	Mean Square	F
Regression	$SSR = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \hat{y}_i)^2$	1	$SSR/1$	MSR/MSE
Residuals	$SSE(R) = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \hat{y}_i)^2$	$n - 2$	$SSE(R)/(n - 2)$	
Lack of Fit	$SS_{LOF} = \sum_{i=1}^m n_i (\bar{y}_i - \hat{y}_i)^2$	$m - 2$	$SS_{LOF}/m - 2$	MS_{LOF}/MS_{PE}
Pure Error	$SS_{PE} = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_i)^2$	$n - m$	$SS_{PE}/n - m$	
Total	$\sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \bar{y})^2$	$n - 1$		

It may be also useful to note that

$$E(SS_{LOF}) = \sigma^2 + \frac{\sum_{i=1}^m n_i (E(y_i) - \beta_0 - \beta_1 x_i)^2}{m - 2}$$

and $E(SS)_{LOF} = \sigma^2$ when we fail to reject the null hypothesis $H_0 : E(y_i) = \beta_0 + \beta_1 x_i$, since the second term becomes 0, and

$$E(SS_{PE}) = \sigma^2$$

3.5 Confidence Intervals in Multiple Regression

3.5.1 Confidence Intervals on Regression Coefficients

We want to construct confidence intervals for the regression coefficients β_j . We use the same assumptions that $\epsilon_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$. So our observations y_i are normally and independently distributed with mean

$$\beta_0 + \sum_{j=1}^k \beta_j x_{ij}$$

and variance σ^2 . $\hat{\beta}$ is a linear combination of the observations, so it follows a normal distribution with a mean vector β and covariance matrix $\sigma^2(X'X)^{-1}$. This implies that the marginal distribution for any $\hat{\beta}_j$ is normal with mean β_j and variance $\sigma^2 C_{jj}$ where C_{jj} is the j th diagonal entry in the matrix $(X'X)^{-1}$. Thus,

$$\frac{\hat{\beta}_j - \beta_j}{\sqrt{\hat{\sigma}^2 C_{jj}}} \sim t_{n-p}$$

So our $100(1 - \alpha)$ percent confidence interval for β_j is

$$\hat{\beta}_j - t_{\alpha/2, n-p} \sqrt{\hat{\sigma}^2 C_{jj}} \leq \beta_j \leq \hat{\beta}_j + t_{\alpha/2, n-p} \sqrt{\hat{\sigma}^2 C_{jj}}$$

Recall that we refer to

$$se(\hat{\beta}_j) = \sqrt{\hat{\sigma}^2 C_{jj}}$$

as the standard error of $\hat{\beta}_j$.

3.5.2 Confidence Intervals On the Mean Response

We can also construct confidence intervals on a specific point, such as $x_{01}, x_{02}, \dots, x_{0k}$. We'll define the vector x_0 as

$$x_0 = \begin{bmatrix} 1 \\ x_{01} \\ x_{02} \\ \vdots \\ x_{0k} \end{bmatrix}$$

The fitted value corresponding to this point is

$$\hat{y}_0 = x_0' \hat{\beta}$$

This is an unbiased estimator for $E(y|x_0)$ since $E(\hat{y}_0) = x_0' \beta = E(y|x_0)$, and the variance is

$$\text{Var}(\hat{y}_0) = \sigma^2 x_0' (X'X)^{-1} x_0$$

Thus our confidence interval for the mean response $E(y|x_0)$ is

$$\hat{y}_0 - t_{\alpha/2, n-p} \sqrt{\hat{\sigma}^2 x_0' (X'X)^{-1} x_0} \leq E(y|x_0) \leq \hat{y}_0 + t_{\alpha/2, n-p} \sqrt{\hat{\sigma}^2 x_0' (X'X)^{-1} x_0}$$

3.5.3 Simultaneous Confidence Intervals on Regression Coefficients

Theorem 3.5.1 (Bonferroni Inequality). *For two events A_1, A_2 , we have that*

$$P(A_1 \cup A_2) = P(A_1) + P(A_2) - P(A_1 \cap A_2) \leq P(A_1) + P(A_2)$$

From DeMorgan's identity, we also have

$$P(A_1^c \cap A_2^c) = 1 - P(A_1 \cup A_2) \geq 1 - P(A_1) - P(A_2)$$

where A_1^c is the complement of A_1 .

If we define the events

$$A_1^c : \hat{\beta}_0 \pm t_{1-\alpha/2, n-2} s(\hat{\beta}_0)$$

$$A_2^c : \hat{\beta}_1 \pm t_{1-\alpha/2, n-2} s(\hat{\beta}_1)$$

where $s(\hat{\beta}_0), s(\hat{\beta}_1)$ are the standard deviations of $\hat{\beta}_0$ and $\hat{\beta}_1$. So the event $(A_1^c \cap A_2^c)$ is the event that the intervals simultaneously cover β_0, β_1 . From Bonferroni's Inequality, if we have $P(A_1) = P(A_2) = \alpha$, then

$$P(A_1^c \cap A_2^c) \geq 1 - P(A_1) - P(A_2) = 1 - 2\alpha$$

In general, if we have p parameters and each confidence interval has confidence, $1 - \frac{\alpha}{p}$, then

$$P\left(\bigcap_{i=1}^p A_i^c\right) \geq 1 - p \frac{\alpha}{p} = 1 - \alpha$$

Conceptually, say we construct $100(1 - \alpha)\%$ confidence intervals on β_1 and β_2 as we have done before with $\alpha = 0.05$. This means, we are 95% confident that the event where β_1 is in its CI occurs, and 95% confident that the event where β_0 is in its interval occurs. So, with the events A_1^c being that β_0 is in its confidence interval, and A_2^c the event β_1 is in its confidence interval, the probability that both β_0 and β_1 are *simultaneously in their respective confidence intervals* is $P(\bigcap A_i^c) \geq 1 - 2\alpha = 0.9$. Thus, if we want to be $100(1 - \alpha)\%$ certain that *both β_0 and β_1 simultaneously lie in their respective confidence intervals*, we need to construct our individual confidence intervals with $100(1 - \alpha/p)\%$ confidence, where p is the number of coefficients.

Chapter 4

Model Adequacy

We've made various assumptions thus far, such as

1. The relationship between the response y and the regressors is linear
2. The error terms ϵ are normally distributed with mean 0 and variance σ^2 .
3. The errors are uncorrelated.

We want to check the validity of these assumptions.

4.1 Residual Analysis

Recall the definition for our observed residuals

$$e_i = y_i - \hat{y}_i$$

Their variance is estimated by

$$\frac{1}{n-p} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{SSE}{n-p} = MSE$$

The residuals are not independent, however, as the n residuals have only $n-p$ degrees of freedom associated with them. This nonindependence of the residuals has little effect on their use for model adequacy checking as long as n is not small relative to the number of parameters p .

4.1.1 Checking for Normality

If the assumption of normality holds, a box plot of the residuals should indicate a symmetric box around the median of 0. A histogram of the residuals can also be used to examine normality. If the residuals follow a similar curve as a normal distribution, then it suggests that the normality assumption is reasonable. The skewness and kurtosis can also help provide insight in this case, a normal distribution has a skewness of 0 and a kurtosis of 3. Finally, a quantile-quantile plot (also known as a qq-plot) compares the quantiles of the residual data with the quantiles from a normal distribution. We compute

$$E_k = \sqrt{MSE} \cdot \Phi^{-1} \left(\frac{k - 0.375}{n + 0.25} \right), \quad k = 1, \dots, n$$

Where Φ is the standard normal, and we plot $e_{(k)}$ vs E_k where $e_{(k)}$ is the residual with rank k . Under normality, we would expect a straight line.

4.1.2 Checking Constant Variance

We use MSE as an estimate for approximating the variance of residual. We can improve the residual scaling by dividing the residuals e_i by the exact standard deviations for the i th residual. Recall that we can use the notion with the hat matrix H to write the vector of residuals as

$$\mathbf{e} = (I - H)\mathbf{Y}$$

where $H = X(X'X)^{-1}X'$ is the hat matrix. Recall the properties of the hat matrix, namely that it is symmetric and idempotent, and $I - H$ has the same properties. We can write the residuals as

$$\mathbf{e} = (I - H)(X\boldsymbol{\beta} + \boldsymbol{\epsilon}) = (I - H)\boldsymbol{\epsilon}$$

The covariance matrix of the residuals is

$$\text{Var}(\mathbf{e}) = \text{Var}[(I - H)\boldsymbol{\epsilon}] = (I - H) \text{Var}(\boldsymbol{\epsilon})(I - H)' = \sigma^2(I - H)$$

This gives us the variance of the i th residual as

$$\text{Var}(e_i) = \sigma^2(1 - h_{ii})$$

and the covariance between residuals e_i and e_j as

$$\text{Cov}(e_i, e_j) = -\sigma^2 h_{ij}$$

Now, we can studentize the residuals to obtain

$$r_i = \frac{e_i}{\sqrt{MSE(1 - h_{ii})}}$$

We can plot the studentized residuals vs fitted values to check for non-constancy of variance. The plot should show a random distribution of points. Conversely, non-constancy of variance would appear as a pattern, an increasing or decreasing collection of points. A scale-location plot can also be used to examine homogeneity, by plotting

$$\sqrt{|r_i|} \text{ vs } \hat{Y}_i$$

If the residuals lie in a narrow band around 0 then there is no evidence to suggest we need corrections. Otherwise, if the residuals show a pattern, either increasing or decreasing, this is a sign that variance is non-constant. If a double-bow pattern appears, this is an indication that the variance in the middle is larger than the variance at the extremes. If the residuals appear to have a quadratic relationship (i.e. a parabola shape), there may be a nonlinear relation that the model has not accounted for.

Chapter 5

Transformations and Weighting to Correct Models

When constructing a regression model, recall that we are making a few assumption

1. The error terms ϵ_i are normally distributed with mean 0 and variance σ^2 , and
2. The error terms are independent and uncorrelated.

In this chapter, the objective is to study methods of building regression models when these assumptions are violated.

5.1 Variance Stabilizing Transformations

The assumption of constant variance is one of the requirements for the regression model. This assumption is commonly violated when the response y has the variance that is functionally related to its mean. For common distributions and functional relationships, we can summarize their useful variance-stabilizing relationships.

Relationship of σ^2 to $E(y)$	Transformation
$\sigma^2 \propto \text{constant}$	$y' = y$ (No transformation)
$\sigma^2 \propto E(y)$	$y' = \sqrt{y}$ (Poisson Data)
$\sigma^2 \propto E(y)[1 - E(y)]$	$y' = \sin^{-1}(\sqrt{y})$ (Binomial Data)
$\sigma^2 \propto E(y)^2$	$y' = \ln(y)$
$\sigma^2 \propto E(y)^3$	$y' = y^{-\frac{1}{2}}$
$\sigma^2 \propto E(y)^4$	$y' = y^{-1}$

In the case of the Poisson distribution, the variance is equal to the mean. So it would be useful to transform the data. If $y \sim \text{Poisson}(\lambda)$, then \sqrt{y} is nearly normally distributed with variance approximately $1/4$ if the mean λ is large.

If we have binomial variable $y \sim \text{Bin}(n, p)$ with mean $m = np$, then we apply the transformation

$$y' = \sin^{-1} \left(\sqrt{\frac{y + c}{n + 2c}} \right)$$

The optimal value of c is $3/8$ when m and $n - m$ are large. The variance is approximately $\frac{1}{4} \left(n + \frac{1}{2} \right)^{-1}$.

5.2 Transformations to Linearize the Model

Another assumption in our regression model is that the relationship between y and the regressors is linear. Sometimes, prior experience or theoretical considerations may indicate that the relationship between y and the regressors is not linear, but may be able to be linearized by using a suitable transformation. These models are known as **intrinsically** or **transformably linear**.

Consider the exponential function,

$$y = \beta_0 e^{\beta_1 x} \epsilon$$

we can transform this model using logarithms to get

$$y' = \ln y = \ln \beta_0 + \beta_1 X + \ln \epsilon = \beta'_0 + \beta_1 x + \epsilon'$$

This transformation required that the new error terms $\epsilon' = \ln \epsilon$ still satisfy our assumptions, namely that they are normally and independently distributed with mean 0 and variance σ^2 .

Various types of reciprocal transformations can also be used, for example

$$y = \beta_0 + \beta_1 \frac{1}{x} + \epsilon$$

This can be linearized using a **reciprocal transformation** for $x' = \frac{1}{x}$ to give us

$$y = \beta_0 + \beta_1 x' + \epsilon$$

Other models can be linearized by reciprocal transformations such as

$$\frac{1}{y} = \beta_0 + \beta_1 x + \epsilon$$

using the transformation $y' = \frac{1}{y}$, and

$$y = \frac{x}{\beta_0 + \beta_1 x}$$

can be linearized with 2 reciprocal transformations. First,

$$y' = \frac{1}{y}$$

then

$$x' = \frac{1}{x}$$

This gives us

$$y' = \beta_0 x' + \beta_1$$

5.2.1 Box-Cox Transformations

Another useful class of transformations when the data appears to be non-normal or non-constant variance is the **power transformation** y^λ , where λ is a parameter to be determined. Box and Cox show how the parameters of the regression model and λ can be estimated simultaneously using the method of maximum likelihood. The appropriate procedure is to use

$$y^{(\lambda)} = \begin{cases} \frac{y^\lambda - 1}{\lambda \dot{y}^{\lambda-1}} & \lambda \neq 0 \\ \dot{y} \ln y & \lambda = 0 \end{cases}$$

Where \dot{y} is the geometric mean of the observations,

$$\dot{y} = \ln^{-1} \left(\frac{1}{n} \sum_{i=1}^n \ln y_i \right)$$

Then we fit the model

$$y^{(\lambda)} = X\beta + \epsilon$$

The divisor $\dot{y}^{\lambda-1}$ turns out to be related to the Jacobian of the transformation converting the response y into $y^{(\lambda)}$. The value of λ is usually determined by trial and error and selecting the value for λ which minimizes the residual sum of squares. We can also construct confidence intervals on λ .

5.3 Generalized and Weighted Least Squares

Linear models that do not satisfy the constant error variance assumption can be fitted by the method of **weighted least squares** to give us constant variances. The idea is to multiply the deviation between the observed and expected value of y_i by a **weight**, w_i , chosen to be inversely proportional to the variance of y_i . In the case of simple linear regression, we'd have

$$Q(\beta_0, \beta_1) = \sum_{i=1}^n w_i \epsilon_i^2 = \sum_{i=1}^n w_i (y_i - \beta_0 - \beta_1 x_i)^2$$

The resulting normal equations from the weighted least squares are

$$\begin{aligned} \hat{\beta}_0 \sum_{i=1}^n w_i + \hat{\beta}_1 \sum_{i=1}^n w_i x_i &= \sum_{i=1}^n w_i y_i \\ \hat{\beta}_0 \sum_{i=1}^n w_i x_i + \hat{\beta}_1 \sum_{i=1}^n w_i x_i^2 &= \sum_{i=1}^n w_i x_i y_i \end{aligned}$$

This model will satisfy that $\text{Var}(\sqrt{w_i} \epsilon_i) = \sigma^2$. We may choose different weights depending on the situation, for example we could choose $w_i = \frac{1}{x_i}$, or $w = \frac{1}{y}$. Another common approach is to preform the usual regression and to estimate the variance with the sample variance s_i^2 for each y_i , then

$$w_i = \frac{1}{s_i^2}$$

5.3.1 Weighted Least Squares

We can define the matrix

$$W = \begin{bmatrix} w_1 & 0 & \cdots & 0 \\ 0 & w_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & w_n \end{bmatrix}$$

Where $\sigma^2 W^{-1}$ is the covariance matrix of ϵ . From the weighted least squares normal equations, we get

$$(X'WX)\hat{\beta} = X'WY$$

This is the multiple linear regression version of the same normal equations given in the simple linear regression model, then we can solve it to get

$$\hat{\beta} = (X'WX)^{-1} X'WY$$

$\hat{\beta}$ is the **weighted least-squares estimator**. If we multiply each of the observed values for the i th observations (including the intercept) by the square root of the weights for the corresponding observations, we transform our data to get

$$X_W = \begin{bmatrix} 1\sqrt{w_1} & x_{11}\sqrt{w_1} & \cdots & x_{1k}\sqrt{w_1} \\ 1\sqrt{w_2} & x_{21}\sqrt{w_2} & \cdots & x_{2k}\sqrt{w_2} \\ \vdots & \vdots & \ddots & \vdots \\ 1\sqrt{w_n} & x_{n1}\sqrt{w_n} & \cdots & x_{nk}\sqrt{w_n} \end{bmatrix}, Y_W = \begin{bmatrix} y_1\sqrt{w_1} \\ y_2\sqrt{w_2} \\ \vdots \\ y_n\sqrt{w_n} \end{bmatrix}$$

The model now becomes

$$Y_W = X_W \beta + \epsilon_W$$

Then the weighted least squares estimate becomes

$$\hat{\beta} = (X_W' X_W)^{-1} X_W' Y_W = (X' W X)^{-1} X' W Y$$

We also define the mean square error as

$$MSE_W = \frac{\sum_{i=1}^n w_i (y_i - \hat{y}_i)^2}{n - p} = \frac{\sum_{i=1}^n w_i e_i^2}{n - p}$$

5.4 Breusch-Pagen Test

The Breusch-Pagen test is a test of constant variance, we assume the model is

$$\log \sigma^2 = \gamma_0 + \gamma_1 X_i$$

We want to test $H_0 : \gamma_1 = 0$. We regress ϵ_i^2 against X_i , and compute the regression sum of squares SSR^* , then our test statistic is

$$\chi^2 = \frac{\frac{SSR^*}{2}}{\frac{SSE}{n}}$$

Under the null hypothesis, $\chi^2 \sim \chi^2(1)$. We reject the null hypothesis when $\chi^2 > \chi^2_{\alpha}(1)$.

Chapter 6

Diagnostics for Leverage and Measures of Influence

A single observation can have a large effect on the regression analysis. We want to detect influential observations that have a large effect on our regression line, and we will be using the hat matrix to do so.

Recall the minimized sum of squares

$$R(\hat{\beta}) = (Y - X\hat{\beta})'(Y - X\hat{\beta}) = Y'(1 - H)Y$$

where $H = X(X'X)^{-1}X'$ is the hat matrix. We define the i^{th} diagonal entry of H as

$$h_{ii} = x_i'(X'X)^{-1}x_i$$

where x_i' is the i^{th} row of X . In the case of simple linear regression with $p = 2$, we have $x_i' = (1, X_i)$ and

$$h_{ii} = \frac{1}{n \sum_{i=1}^n X_i^2 - (\sum_{i=1}^n X_i)^2} x_i' \begin{pmatrix} \sum_{i=1}^n X_i^2 & -\sum_{i=1}^n X_i \\ -\sum_{i=1}^n X_i & n \end{pmatrix} x_i = \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

6.1 Leverage

Definition 6.1.1. The *leverage* of the i^{th} observation is defined as h_{ii} .

We can gain further insight by writing the mean \bar{X} in terms of the mean with the i^{th} observation removed, which we denote $\bar{X}_{(i)}$.

$$\bar{X} = \frac{1}{n}(X_i + (n-1)\bar{X}_{(i)})$$

Then,

$$X_i - \bar{X} = X_i - \frac{1}{n}(X_i + (n-1)\bar{X}_{(i)}) = \frac{n-1}{n}(X_i - \bar{X}_{(i)})$$

Then, we can rewrite the leverage of the i^{th} observation as

$$h_{ii} = \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{1}{n} + \left(\frac{n-1}{n}\right)^2 \frac{(X_i - \bar{X}_{(i)})^2}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

This tells us that the leverage of the i^{th} observation will be large if X_i is far from the mean of the observations.

6.1.1 Properties of Leverage

The leverage can be used to detect influential observations, from the fact that

$$\text{Trace}(H) = \text{Trace}(X(X'X)^{-1}X') = \text{Trace}((X'X)^{-1}X'X) = \text{Trace}(I_p) = p$$

So the average leverage is $\sum_{i=1}^n h_{ii}/n = p/n$. If $h_{ii} > 2p/n$, then the i^{th} observation is considered to be influential. Note that not all high leverage observations are influential.

6.2 Measures of Influence

6.2.1 Difference in Fit

We can measure the difference in fit by comparing the sum of squares with and without the i^{th} observation. We define the **Difference in fit** as

$$DFFITs_i = \frac{\hat{Y}_i - \hat{Y}_{i(i)}}{\sqrt{MSE_{(i)} h_{ii}}} = t_i \left(\frac{h_{ii}}{1 - h_{ii}} \right)^{\frac{1}{2}}$$

where t_i represents the Studentized deleted residual

$$t_i = e_i \left(\frac{n - p - 1}{SSE(1 - h_{ii}) - e_i^2} \right)^{\frac{1}{2}}$$

This shows that we can calculate the difference in fit using the error sum of squares and the hat matrix. The value for $DFFITs_i$ represents the number of estimated standard deviations of \hat{Y}_i that the fitted value increases or decreases with the inclusion of the i^{th} observation. If the i^{th} observation is an outlier and has high leverage, then

$$\left(\frac{h_{ii}}{1 - h_{ii}} \right)^{\frac{1}{2}} > 1$$

and $DFFITs$ will be large in absolute value. Influential cases are flagged for small datasets if

$$|DFFITs_i| > 1$$

For large datasets, we use

$$|DFFITs_i| > 2\sqrt{\frac{p}{n}}$$

6.2.2 Cook's Distance

Cook's distance considers the influence of the i^{th} observation on the entire regression line. We define Cook's distance as

$$D_i = \frac{\sum_{j=i}^n (\hat{Y}_j - \hat{Y}_{j(i)})^2}{pMSE} = \frac{e_i^2}{pMSE} \left(\frac{h_{ii}}{(1 - h_{ii})^2} \right)$$

Cook's distance is large if the residual is large and the leverage is moderate, or if the residual is moderate and the leverage is large, or if both are large. Influential cases are flagged when $D_i > 1$.

6.2.3 Difference in Coefficients

The **DFBETAS** are the differences in the estimated regression coefficients with and without the i^{th} observation. We define the DFBETAS as

$$DFBETAS_{(i)} = \frac{\hat{\beta}_k - \hat{\beta}_{k(i)}}{\sqrt{MSE_{(i)} c_{ii}}}$$

where c_{ii} is the i^{th} diagonal entry of $(X'X)^{-1}$. A large value of $DFBETAS_{(i)}$ indicates that the i^{th} observation has a large effect on the k^{th} regression coefficient. Influential cases are flagged if

$$DFBETAS_{(i)} > \begin{cases} 1 & \text{for small datasets} \\ 2/\sqrt{n} & \text{for large datasets} \end{cases}$$

Chapter 7

Polynomial Regression Models

7.1 Polynomial Models in One Variable

The k -order polynomial regression model in one variable is defined as

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \cdots + \beta_k X^k + \epsilon$$

It is important to keep the order k as small as possible. A lower order model in a transformed variable is almost always preferable to a higher order model in the original variable.

Often, orthogonal polynomials are used in modeling because they simplify the fitting process.

$$Y_i = \beta_0 P_0(X_i) + \beta_1 P_1(X_i) + \beta_2 P_2(X_i) + \cdots + \beta_k P_k(X_i) + \epsilon_i$$

where P_j is a j order polynomial satisfying

$$\sum_{i=1}^n P_j(X_i) P_l(X_i) = 0, \quad j \neq l$$
$$P_0(X_i) = 1$$

The least squares estimates are given by

$$\hat{\beta}_j = \frac{\sum_{i=1}^n P_j(X_i) Y_i}{\sum_{i=1}^n P_j^2(X_i)}, \quad j = 0, 1, \dots, k$$

The model then becomes $y = X\beta + \epsilon$ where X is

$$X = \begin{bmatrix} P_0(X_1) & P_1(X_1) & \cdots & P_k(X_1) \\ P_0(X_2) & P_1(X_2) & \cdots & P_k(X_2) \\ \vdots & \vdots & \ddots & \vdots \\ P_0(X_n) & P_1(X_n) & \cdots & P_k(X_n) \end{bmatrix}$$

Since this matrix has orthogonal columns, $X'X$ becomes

$$X'X = \begin{bmatrix} \sum_{i=1}^n P_0^2(X_i) & 0 & \cdots & 0 \\ 0 & \sum_{i=1}^n P_1^2(X_i) & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sum_{i=1}^n P_k^2(X_i) \end{bmatrix}$$

The residual sum of squares is

$$SSE(k) = SST - \sum_{j=1}^k \hat{\beta}_j \left(\sum_{i=1}^n P_j(X_i) Y_i \right)$$

and the regression sum of squares is

$$SSR(\beta_j) = \hat{\beta}_j \sum_{i=1}^n P_j(X_i)Y_i$$

To test significance of the highest order term, we should test $H_0 : \beta_k = 0$, using the F statistic

$$F_0 = \frac{SSR(\beta_k)}{SSE(k)/(n-k-1)} = \frac{\hat{\beta}_k \sum_{i=1}^n P_k(X_i)Y_i}{SSE(k)/(n-k-1)}$$

The advantage of using orthogonal polynomials is that the model can be fitted sequentially. Models can be done by computers so this is not as important as it once was.

7.2 Polynomial Models in Two Variables

When 2 or more variables are involved, cross terms are included in the model. For 2 variables, the model is

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_{11} X_1^2 + \beta_{22} X_2^2 + \beta_{12} X_1 X_2 + \epsilon$$

These are referred to as response surfaces.

7.3 Indicator Regression Models

When working with qualitative or categorical data, indicator functions can be used as "dummy" variables.

$$X_2 = \begin{cases} 1 & \text{Male} \\ 0 & \text{Female} \end{cases}$$

An application of this is where if you want to fit a simple model as a function of gender, we may have

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$$

With X_2 being the gender variable, so

$$Y = \begin{cases} \beta_0 + \beta_1 X_1 + \beta_2 + \epsilon & \text{Male} \\ \beta_0 + \beta_1 X_1 + \epsilon & \text{Female} \end{cases}$$

Chapter 8

Multicollinearity

8.1 Properties of Multicollinearity

Multicollinearity refers to the situation where two or more of the predictor variables exhibit a near-linear relationship. If there is no linear relationship between the regressors, they are said to be **orthogonal**. The symptoms of multicollinearity are:

1. Large variation in the estimated coefficients when a new variable is either added or deleted.
2. Non-significant results in individual tests on the coefficients of important variables.
3. Large coefficients of simple correlation between pairs of variables.
4. Wide confidence interval for the regression coefficients of important variables.

The main difficulty is that the matrix $(X'X)$ may not be invertible. Multicollinearity affects the interpretation of the coefficients in that they may vary in value.

8.2 Detecting Multicollinearity

Consider the case of two predictor variables X_1 and X_2 . If the variables are standardized, then the matrix becomes

$$(X'X) = \frac{1}{1 - r_{12}^2} \begin{pmatrix} 1 & r_{12} \\ r_{12} & 1 \end{pmatrix}$$

where r_{12} is the correlation between the two variables. The variance covariance matrix of the coefficients is given by

$$\sigma^2(X'X)^{-1} = \sigma^2 \frac{1}{1 - r_{12}^2} \begin{pmatrix} 1 & -r_{12} \\ -r_{12} & 1 \end{pmatrix}$$

As $|r_{12}| \rightarrow 1$, the variance $\text{Var}(\hat{\beta}_k) \rightarrow \infty$, and the covariance $\text{Cov}(\hat{\beta}_1, \hat{\beta}_2) \rightarrow \pm\infty$ depending on if $r_{12} \rightarrow +1$ or $r_{12} \rightarrow -1$. The estimates are

$$\hat{\beta} = (X'X)^{-1}X'Y$$

which can be written as the individual estimates

$$\hat{\beta}_1 = \frac{r_{1Y} - r_{12}}{1 - r_{12}^2}, \quad \hat{\beta}_2 = \frac{r_{2Y} - r_{12}}{1 - r_{12}^2}$$

In general, the diagonal elements of $(X'X)^{-1}$ are $C_{jj} = \frac{1}{1 - R_j^2}$ where R_j^2 is the R -square value obtained from the regression of X_j on the other $p - 1$ variables. If there is a strong multicollinearity between X_j and the other $p - 1$ variables, then

$$R_j^2 \approx 1, \quad \text{Var}(\hat{\beta}_j) = \frac{\sigma^2}{1 - R_j^2} \rightarrow \infty$$

Multicollinearity tends to produce least squares estimates $\hat{\beta}_j$ that are too far large in absolute value. To see this, consider the squared distance from $\hat{\beta}$ and β :

$$L^2 = \|\hat{\beta} - \beta\|^2 = (\hat{\beta} - \beta)'(\hat{\beta} - \beta)$$

The expected value for the squared distance is

$$\begin{aligned} E(L^2) &= E(\hat{\beta} - \beta)'(\hat{\beta} - \beta) \\ &= \sum_{j=1}^p E(\hat{\beta}_j - \beta_j)^2 \\ &= \sum_{j=1}^p \text{Var}(\hat{\beta}_j) = \sigma^2 \text{Trace}(X'X)^{-1} \\ &= \sigma^2 \sum_{j=1}^p \frac{1}{\lambda_j} \end{aligned}$$

where λ_j are the eigenvalues of $X'X$. If the matrix $X'X$ is ill-conditioned because of multicollinearity, at least one of the eigenvalues λ_j will be small, which will cause the distance from the least squares estimate $\hat{\beta}$ and the true parameters β to be large. Using this result, we can show that

$$E(\hat{\beta} - \beta)'(\hat{\beta} - \beta) = E(\hat{\beta}'\hat{\beta} - 2\hat{\beta}'\beta + \beta'\beta) = \sigma^2 \sum_{j=1}^n \frac{1}{\lambda_j}$$

This gives us

$$E(\|\hat{\beta}\|^2) = E(\hat{\beta}'\hat{\beta}) = \beta'\beta + \sigma^2 \text{Trace}(X'X)^{-1}$$

Which means, the vector $\hat{\beta}$ is generally longer than the vector β . This implies that the method of least squares produces estimated regression coefficients that are too large in absolute value.

8.2.1 Variance Inflation Factors

Suppose the regression model is fitted using the standardized predictor variables,

$$\begin{aligned} Y_i^* &= \frac{1}{\sqrt{n-1}} \left(\frac{Y_i - \bar{Y}}{s_Y} \right) \\ X_{ik}^* &= \frac{1}{\sqrt{n-1}} \left(\frac{X_{ik} - \bar{X}_k}{s_k} \right), \quad k = 1, \dots, p-1 \end{aligned}$$

The model then becomes

$$\begin{aligned} Y_i^* &= \sum_{k=1}^{p-1} \beta_k^* X_{ik}^* + \epsilon_i^* \\ \beta_k &= \left(\frac{s_Y}{s_k} \right) \beta_k^* \\ \beta_0 &= \bar{Y} - \sum_{i=1}^{k-1} \beta_i \bar{X}_i \\ r_{XX} \hat{\beta}^* &= r_{YX} \end{aligned}$$

where r_{XX} is the correlation matrix

$$r_{XX} = \begin{pmatrix} 1 & r_{12} & r_{13} & \cdots & r_{1k} \\ r_{12} & 1 & r_{23} & \cdots & r_{2k} \\ r_{13} & r_{23} & 1 & \cdots & r_{3k} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ r_{1k} & r_{2k} & r_{3k} & \cdots & 1 \end{pmatrix}$$

and r'_{YX} is the vector of simple correlation between x_j and the response y ,

$$r_{jy} = \text{Cor}(Y, X_j)$$

The variance of the estimates β is

$$\text{Var}(\beta) = \sigma^2 r_{XX}^{-1}$$

We define the variance inflation factor as

$$\text{VIF}_j = C_{jj} = (1 - R_j^2)^{-1}$$

where R_k^2 is the coefficient of multiple determination when X_k is regressed on the $p - 2$ other X variables. So

$$\text{Var}(\hat{\beta}_k) = \sigma^2 (1 - R_k^2)^{-1}$$

The variance inflation factor is 1 when the coefficient of determination is 0, in other words whenever X_k is not linearly related to the other X variables in the model. As a general rule, a value of $\text{VIF}_j > 10$ indicates multicollinearity exists. We may also compute the average of the variance inflation factors,

$$\overline{\text{VIF}} = \frac{\sum_{k=1}^{p-1} \text{VIF}_k}{p - 1}$$

A mean value greater than 1 indicates serious multicollinearity.

8.3 Ridge Regression

Ridge regression is considered as a remedial measure for multicollinearity. We transform the normal equation

$$(X'X)\hat{\beta} = X'Y$$

using the standardized variables so that it becomes

$$r_{XX}\hat{\beta} = r_{YX}$$

The eigenvalues can also be used to measure the extent of multicollinearity. If one or more eigenvalues are small, then there are near linear dependences in the columns of $X'X$. We define the condition number κ and condition indices κ_j of $X'X$ as

$$\kappa = \frac{\lambda_{\max}}{\lambda_{\min}}, \quad \kappa_j = \frac{\lambda_{\max}}{\lambda_j}$$

If $\kappa < 100$, then there is no serious multicollinearity. If $100 < \kappa < 1000$, then there is moderate to strong multicollinearity. If $\kappa > 1000$, then there is severe multicollinearity. We define the ridge estimator $\hat{\beta}_R$ as the solution to the equation

$$(r_{XX} + cI)\hat{\beta}_R = r_{YX}$$

where $c \geq 0$ is a constant. Solving this yields

$$\hat{\beta}_R = (r_{XX} + cI)^{-1} r_{YX}$$

The constant c reflects the fact that the ridge estimators will be biased but they tend to be more stable or less variable than the ordinary least squares estimators. Ridge regression can also be obtained from the method of penalized regression, with the following system of equations

$$\begin{aligned} r_{Y1} &= (1 + c)\hat{\beta}_{R1} + r_{12}\hat{\beta}_{R2} + \cdots + r_{1,p-1}\hat{\beta}_{R,p-1} \\ r_{Y2} &= r_{21}\hat{\beta}_{R1} + (1 + c)\hat{\beta}_{R2} + \cdots + r_{2,p-1}\hat{\beta}_{R,p-1} \\ r_{Y,p-1} &= r_{p-1,1}\hat{\beta}_{R1} + r_{p-1,2}\hat{\beta}_{R2} + \cdots + (1 + c)\hat{\beta}_{R,p-1} \end{aligned}$$

We can write the penalized least squares as

$$Q = \sum (Y_i - \beta_1 X_{i1} - \cdots - \beta_{p-1} X_{i,p-1})^2 + c \sum_{j=1}^{p-1} \beta_j^2$$

Differentiating with respect to each of the parameters gives us the equations above.

8.3.1 LASSO

Another approach to ridge regression is to use a penalty function

$$c \sum_{j=1}^{p-1} |\beta_j|$$

which permits some regression coefficients to be 0. LASSO stands for Least Absolute Shrinkage and Selection Operator

Chapter 9

Building the Regression Model

In application, when dealing with a large number of variables, we want to be able to select the best subset of variables to include in our model. Finding this appropriate subset of variables is often called the **variable selection problem**. If we have $p-1$ predictors available, then there are $2^p - 1$ possible models we can construct. One approach is to select **all possible regression** and test $2^p - 1$ possible models.

9.1 Criteria for Model Selection

We may choose our model based on the largest value of explained variation R^2 . A plot of the mean square residual vs the number of variables will appear as a parabola, from which we can determine the best fitting model. An adjusted R^2 takes into account the values of n and p ,

$$R_{a,p}^2 = 1 - \left(\frac{n-1}{n-p} \right) \frac{SSE}{SSTO} = 1 - \frac{MSE(p)}{SSTO/(n-1)}$$

We want to minimize $MSE(p)$ which in turn maximizes the adjusted R^2 .

9.1.1 Mallows's C_p

Mallows proposed a criterion that is related to the mean square error of a fitted value,

$$E(\hat{Y}_i - E(Y_i))^2 = (E(Y_i) - E(\hat{Y}_i))^2 + \text{Var}(\hat{Y}_i)$$

Suppose that the true model has q predictor variables,

$$Y = X_q \beta_Q + \epsilon$$

and that we fit a model with p predictors. Let H_p be the hat matrix using only p variables. We define the bias for the i^{th} fitted value as

$$E(\hat{Y}_i) - \mu_i$$

where μ_i is the true mean $E(Y_i)$. Let the total squares bias for the p -term equation be

$$SSB(p) = \sum_{i=1}^n (E(Y_i) - E(\hat{Y}_i))^2$$

Then we define the standardized total mean square error as

$$\Gamma_p = \frac{1}{\sigma^2} \left(\sum_{i=1}^n (E(\hat{Y}_i) - \mu_i)^2 + \sum_{i=1}^n \text{Var}(\hat{Y}_i) \right) = \frac{SSB(p)}{\sigma^2} + \frac{1}{\sigma^2} \sum_{i=1}^n \text{Var}(\hat{Y}_i)$$

It can be shown that

$$\sum_{i=1}^n \text{Var}(\hat{Y}_i) = p\sigma^2$$

and the expected value of the residual sum of squares is

$$E(SSE_p) = SSB_p + (n-p)\sigma^2$$

Thus,

$$\Gamma_p = \frac{SSB_p}{\sigma^2} + \frac{1}{\sigma^2} \sum_{i=1}^n \text{Var}(\hat{Y}_i) = \frac{E(SSE_p) - (n-p)\sigma^2}{\sigma^2} + p = \frac{E(SSE_p)}{\sigma^2} - n + 2p$$

The vector of residuals becomes

$$\hat{\epsilon}_p = (1 - H_p)Y$$

and the sum of squares is

$$SSE_p = \hat{\epsilon}_p' \hat{\epsilon}_p$$

So,

$$\text{Bias} = E(\hat{\epsilon}) = (1 - H_p)E(Y) = E(Y) - E(\hat{Y})$$

When $p = q$, we have that the bias is equal to 0. Now using the fact that $(1 - H_p)$ is idempotent, we have that

$$\begin{aligned} E(SSE_p) &= E(\hat{\epsilon}_p' \hat{\epsilon}_p) \\ &= E(Y'(1 - H_p)Y) \\ &= E(Y'(1 - H_p)(1 - H_p)Y) \\ &= \sigma^2 \text{Trace}(I - H_p) + \text{Bias}' \text{Bias} \\ &= \sigma^2(n - p) + \sum_{i=1}^n (E(Y_i) - E(\hat{Y}_i))^2 \end{aligned}$$

We can show that the sum of the variances of the fitted values is $p\sigma^2$,

$$\begin{aligned} \sum_{i=1}^n \text{Var}(\hat{Y}_i) &= \sigma^2 \sum_{i=1}^n (x_i'(X'X)^{-1}x_i) \\ &= \sigma^2 \text{Trace}(X'(X'X)^{-1}X) \\ &= \sigma^2 \text{Trace}(H) \\ &= p\sigma^2 \end{aligned}$$

To reiterate, this gives us our standardized total mean square error as

$$\Gamma_p = \frac{SSE_p}{\sigma^2} - n + 2p$$

When we use MSE to estimate σ^2 , and SSE_p so estimate $E(SSE_p)$, we get Mallows's C_p statistic,

$$C_p = \frac{SSE_p}{MSE} - n + 2p$$

If there is negligible bias in the p -term model, then $E(SSE_p) \approx (n - p)\sigma^2$ and $C_p \approx p$. Mallows technique to find the optimal subset of variables is to plot C_p vs p for all the possible regression, and models with small bias will be close to the line $C_p = p$. Models with large bias will be above the line, and values below the line are considered to have no bias.

9.1.2 Akaike Information Criterion

Akaike proposed an information criterion (AIC) that is based on the maximized the expected entropy of the model. Essentially, AIC is a penalized log-likelihood function. Let L be the likelihood function for a specific model, AIC is defined as

$$\text{AIC} = -2 \ln L + 2p$$

In the case of ordinary least squares regression, it becomes

$$\text{AIC}_p = n \ln(SSE_p) - n \ln n + 2p$$

There are several Bayesian extension of AIC, such as the Shwartz's Bayesian criterion,

$$\text{BIC}_{\text{Sch}} = n \ln(SSE_p) - n \ln n + p \ln n$$

This criterion places a larger penalty on adding regressors as the sample size increases and is the one used in R.

9.1.3 Prediction Sum of Squares Criterion

Frequently, regression models are used for prediction of future observations or estimation of the mean response. Generally, we want to select the regressors that minimize the prediction error. So, we select the model that minimizes the prediction sum of squares,

$$\text{PRESS}_p = \sum_{i=1}^n (Y_i - \hat{Y}_{(i)})^2 = \sum_{i=1}^n \left(\frac{e_i}{1 - h_{ii}} \right)^2$$

9.2 Computational Techniques for Variable Selection

Evaluating model selection criteria for all possible models requires a large amount of computations. We look at several techniques to reduce the computational burden.

9.2.1 Forward Selection

The procedure begins with the assumption that there are no regressors in the model other than the intercept. Then, we insert regressors one at a time. The first regressor selected to be added in the model is the one with the largest simple correlation, say x_1 , with the response y . This is also the variable with the greatest t -statistic in absolute value. A preselected F -value, which we call F_{IN} , is used to determine whether or not to add a regressor to the model.

Then, after selecting the first variable, we choose the second variable which has the largest correlation with the response after adjusting for the effect of the first regressor. These correlations are called **partial correlation**. They are computed between the residuals from the first regression and the residuals from the regressions of the other candidate regressors on x_1 . That is, the residuals from

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1$$

and the residuals on the other candidate regressors,

$$\hat{x}_j = \hat{\alpha}_{0j} + \hat{\alpha}_{1j} x_1, \quad j = 2, \dots, K$$

Once we select x_2 , this implies that the largest partial F -statistic is

$$F = \frac{SSR(x_2|x_1)}{MSE(x_1, x_2)}$$

If $F > F_{\text{IN}}$, then we add x_2 to the model. We then check the t -value and if it is below a preset limit, we may need to drop a variable from the model. We repeat this procedure until the largest partial F -statistic is less than F_{IN} , or until we have added all the variables.

9.2.2 Backward Elimination

Backward selection works in the opposite direction. We begin with a model that includes all K candidate regressors. We then compute the partial F -statistic for each regressor as if it were the last one to enter the model. The smallest of these partial F -statistics is compared with a preselected F -value, F_{OUT} . If the smallest partial F -value is less than F_{OUT} , that regressor is removed from the model, and we fit the regression model with the remaining $K - 1$ regressors. The new partial F -statistic is calculated, and we repeat this process. The procedure stops when the smallest partial F value is not less than the preselected cutoff value, F_{OUT} .

9.2.3 Stepwise Regression

Stepwise regression is a modification of forward selection. In each step, all regressors entered into the model thus far are reassessed with their partial F statistics. A regressor added at an earlier step may now be redundant because of the relationship it has with the new regressors. If the partial F statistic for a variable is less than F_{OUT} , then that variable is removed from the model. This process requires 2 cutoff values, F_{IN} and F_{OUT} . Some may choose $F_{\text{IN}} = F_{\text{OUT}}$, however $F_{\text{IN}} > F_{\text{OUT}}$ is more common since it makes the criteria for adding variables more strict than the criteria for removing variables.

Chapter 10

Logistic Regression and General Linear Models

Sometimes we wish to model a binary response variable, a variable that can only take on two values say, 0 or 1. Consider the model where we estimate the likelihood that a person is wearing a life jacket,

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

where

$$Y_i = \begin{cases} 1 & \text{with probability } \pi_i \\ 0 & \text{with probability } 1 - \pi_i \end{cases}$$

Then,

$$E(Y_i) = \pi_i$$

The least squares model is problematic because

- The variance $\text{Var}(Y_i) = \pi_i(1 - \pi_i)$ is not constant.
- The error terms ϵ_i are not normally distributed.
- A linear model cannot guarantee that \hat{Y}_i is between 0 and 1.

10.1 Logistic Distribution

The logistic distribution has the density function

$$f(x) = \frac{e^x}{(1 + e^x)^2}$$

with a cumulative distribution function

$$F(t) = \frac{e^t}{1 + e^t}$$

It can be shown that $E(X) = 0$ and $\text{Var}(X) = \frac{\pi^2}{3}$. Suppose a random variable Y is binary with

$$Y_i = \begin{cases} 1 & \beta_0^* + \beta_1^* X_i + \epsilon_i^* < k \\ 0 & \beta_0^* + \beta_1^* X_i + \epsilon_i^* > k \end{cases}$$

for some constant k , and ϵ_i^* has a logistic distribution. Then,

$$\begin{aligned} \pi_i &= P(Y_i = 1) = P(\beta_0^* + \beta_1^* X_i + \epsilon_i^* < k) \\ &= F(k - \beta_0^* - \beta_1^* X_i) \\ &= F(\beta_0 + \beta_1 X_i) \\ &= \frac{e^{\beta_0 + \beta_1 X_i}}{1 + e^{\beta_0 + \beta_1 X_i}} \end{aligned}$$

where $\beta_0 = k - \beta_0^*$ and $\beta_1 = -\beta_1^*$.

10.2 Estimating Parameters

We estimate parameters based on maximizing the likelihood, so it is useful to note that

$$\ln\left(\frac{\pi_i}{1-\pi_i}\right) = \beta_0 + \beta_1 X_i$$

Then,

$$\begin{aligned} \prod_{i=1}^n f(y_i) &= \prod_{i=1}^n \pi_i^{y_i} (1-\pi_i)^{1-y_i} \\ &= \prod_{i=1}^n \left[\left(\frac{\pi_i}{1-\pi_i} \right)^{y_i} (1-\pi_i) \right] \\ \Rightarrow \ln\left(\prod_{i=1}^n f(y_i)\right) &= \sum_{i=1}^n y_i \ln\left(\frac{\pi_i}{1-\pi_i}\right) + \sum_{i=1}^n \ln(1-\pi_i) \\ &= \sum_{i=1}^n y_i (\beta_0 + \beta_1 X_i) - \sum_{i=1}^n \ln(1 + e^{\beta_0 + \beta_1 X_i}) \end{aligned}$$

There is no closed form solution to the maximum likelihood estimator, so we use numerical methods to obtain $\hat{\beta}_0$, and $\hat{\beta}_1$. So we get

$$\hat{\pi} = \frac{\exp(\hat{\beta}_0 + \hat{\beta}_1 X_i)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}_1 X_i)}$$

10.2.1 Interpretation of the Parameters

Consider the fitted value at a specific value of X , say X_0 . We will define the linear predictor as

$$\eta = \ln \frac{\pi_i}{1-\pi_i} = \beta_0 + \beta_1 X_i$$

Then, the linear predictor at X_0 is

$$\hat{\eta}(X_0) = \hat{\beta}_0 + \hat{\beta}_1 X_0$$

Then similarly for $X_0 + 1$,

$$\hat{\eta}(X_0 + 1) = \hat{\beta}_0 + \hat{\beta}_1 (X_0 + 1)$$

We defined η to be the log-odds, which is simply the logarithm of the odds of event occurring. The odds are the probability an event occurs divided by the probability it does not occur. So, in this case our probability is π_i , so the difference in the log odds is

$$\hat{\eta}(X_0 + 1) - \hat{\eta}(X_0) = \hat{\beta}_1$$

This gives us

$$\hat{\eta}(X_0 + 1) - \hat{\eta}(X_0) = \ln(\text{odds}_{X_0+1}) - \ln(\text{odds}_{X_0}) = \ln\left(\frac{\text{odds}_{X_0+1}}{\text{odds}_{X_0}}\right)$$

Computing the anti-log of both sides gives us

$$\hat{O}_R = \frac{\text{odds}_{X_0+1}}{\text{odds}_{X_0}} = e^{\hat{\beta}_1}$$

This is known as the odds ratio. The odds ratio is the estimated increase in the probability of successes associated with a one unit change in the value of the predictor variable X . For a change of d units, then

$$\hat{O}_R = e^{d\hat{\beta}_1}$$

10.2.2 Repeat Observations

Suppose that we have repeat observations at each of the levels of X , and set Y_i to be the number of 1's observed for the i th observation. Let n_i be the number of trials at each observation. Then $Y_i \sim \text{Bin}(n_i, \pi_i)$, and the likelihood is

$$L(\beta_0, \beta_1) = \prod_{i=1}^n \binom{n_i}{Y_i} \pi_i^{Y_i} (1-\pi_i)^{n_i-Y_i}$$

We want to maximize this, so

$$\log L(\beta_0, \beta_1) = \sum_{i=1}^n \left(\log \binom{n_i}{Y_i} + Y_i \log \pi_i + (n - Y_i) \log(1 - \pi_i) \right)$$

10.3 Multiple Logistic Regression

We can fit a multiple logistic regression model by replacing $\beta_0 + \beta_1 X$ with

$$X'_i \beta = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_{i,p-1} X_{i,p-1}$$

$$E(Y) = \frac{\beta' X}{1 + \beta' X}$$

$$\log \frac{\pi}{1 - \pi} = \beta' X$$

10.3.1 Inference on Model Parameters

The maximum likelihood estimators are asymptotically normal with variances and covariances that are functions of the second order partial derivatives of the likelihood function. Let

$$G = \left(\frac{\partial^2 L(\beta)}{\partial \beta_i \partial \beta_j} \right)$$

be the Hessian were

$$\log L(\beta) = \sum_{i=1}^n Y_i (X'_i \beta) - \sum_{i=1}^n \log(1 + e^{X'_i \beta})$$

The linear predictor is $X'_i \hat{\beta}$ and the fitted values are

$$\hat{Y}_i = \hat{\pi}_i = \frac{\exp(X'_i \hat{\beta})}{1 + \exp(X'_i \hat{\beta})}$$

It can be shown that $\hat{\beta}$ is an unbiased estimator, and the variance is

$$\text{Var}(\hat{\beta}) = (X' V X)^{-1}$$

Where $V = \text{diag}(n \hat{\pi}_i (1 - \hat{\pi}_i))$. Furthermore, we have

$$\frac{\hat{\beta}_k - \beta_k}{s(\hat{\beta}_k)} \sim N(0, 1), \quad k = 0, \dots, p-1$$

We use this for testing and constructing confidence intervals. The confidence interval on the parameters is given by

$$\hat{\beta}_k - z_{\alpha/2} s(\hat{\beta}_k) \leq \hat{\beta}_k \leq \hat{\beta}_k + z_{\alpha/2} s(\hat{\beta}_k)$$

For the odds ratio,

$$\exp(\hat{\beta}_1 - z_{\alpha/2} s(\hat{\beta}_1)) \leq \hat{O}_R \leq \exp(\hat{\beta}_1 + z_{\alpha/2} s(\hat{\beta}_1))$$

10.4 Testing

To test whether several coefficients are 0, we use the likelihood ratio test where we compare the full model (FM) with the reduced model (RM). We define

$$G^2 = -2 \log \left(\frac{L(RM)}{L(FM)} \right) = 2 \log \left(\frac{L(FM)}{L(RM)} \right)$$

which is asymptotically χ^2 with degrees of freedom equal to the difference in the number of parameters in the two models if the null hypothesis is true. We reject for large values of G^2 , when $G^2 > \chi^2_{p-q}$.

In the case of simple logistic regression, the full model is the one that has been fitted, whereas the reduced model is the one with constant probability of success,

$$E(Y) = \pi = \frac{e^{\beta_0}}{1 + e^{\beta_0}}$$

The maximum likelihood estimator of π in the reduced model is $\hat{\pi} = \frac{Y}{n}$, so,

$$\ln L(RM) = y \ln y + (n - y) \ln(n - y) - n \ln n$$

Thus, the likelihood ratio statistic for testing significance of regression is

$$L = 2(\ln L(FM) - \ln L(RM)) = 2 \left(\sum_{i=1}^n y_i \ln \hat{\pi}_i + \sum_{i=1}^n (n_i - y_i) \ln(1 - \hat{\pi}_i) - y \ln y - (n - y) \ln(n - y) + n \ln n \right)$$

10.4.1 Test for Goodness of Fit

We would like to test

$$H_0 : E(Y) = \left(1 + e^{-X'\beta}\right)^{-1}$$

$$H_1 : E(Y) \neq \left(1 + e^{-X'\beta}\right)^{-1}$$

We will make use of a Pearson chi-square goodness of fit test. The expected number of successes is $n_i \hat{\pi}_i$, and the expected number of failures is $n_i(1 - \hat{\pi}_i)$. Then, the Pearson chi-square statistic is

$$\chi^2 = \sum_{i=1}^n \left(\frac{(Y_i - n_i \hat{\pi}_i)^2}{n_i \hat{\pi}_i} + \frac{(n_i - Y_i - n_i(1 - \hat{\pi}_i))^2}{n_i(1 - \hat{\pi}_i)} \right) = \sum_{i=1}^n \frac{y_i - n_i \hat{\pi}_i}{n_i \hat{\pi}_i(1 - \hat{\pi}_i)}$$

We reject the null hypothesis when

$$\chi^2 > \chi_{\alpha, n-p}^2$$

10.4.2 Hosmer-Lemeshow Test

When there are no replicates on the regressor variables, the observations can be grouped to perform a goodness-of-fit test known as the Hosmer-Lemeshow test. In this procedure, we classify the observations into g groups, generally $g = 10$ is used. The observed number of successes O_j and failures $N_j - O_j$ are compared with the expected frequency of each group, $N_j \bar{\pi}_j$ where N_j is the number of observations in the j th group. The average estimated success probability in the j th group is

$$\bar{\pi}_j = \sum_{i \in \text{group } j} \frac{\hat{\pi}_i}{N_j}$$

The Hosmer-Lemeshow statistic is just a Pearson chi-square goodness-of-fit statistic comparing observed and expected frequencies, and is given as

$$HL = \sum_{j=1}^g \frac{(O_j - N_j \bar{\pi}_j)^2}{N_j \bar{\pi}_j} + \sum_{j=1}^g \frac{(N_j - O_j - N_j(1 - \bar{\pi}_j))^2}{N_j(1 - \bar{\pi}_j)} = \sum_{j=1}^g \frac{(O_j - N_j \bar{\pi}_j)^2}{N_j \bar{\pi}_j(1 - \bar{\pi}_j)}$$

If the fitted model is correct, the HL statistic follows a chi-square distribution with $g - 1$ degrees of freedom. We reject the null hypothesis for large values of HL.

10.4.3 Deviance of Goodness of Fit Test

Here we compare the current model to a saturated model where each observation (or group when $n_i > 1$) has its own probability of success estimated by Y_i/n_i . Under the reduced model

$$E(Y_i) = \left(1 + e^{-X_i'\beta}\right)^{-1}$$

and under the saturated model (also called the full model)

$$E(Y_i) = \pi_i$$

The deviance goodness of fit statistic is

$$\text{Dev}(X_0, X_1, \dots, X_{p-1}) = -2(\ln L(RM) - \ln L(FM)) = -2 \sum_{i=1}^n \left(Y_i \ln \left(\frac{Y_i}{n_i \hat{\pi}_i} \right) \right) - 2 \sum_{i=1}^n \left((n_i - Y_i) \ln \left(\frac{n_i - Y_i}{n_i(1 - \hat{\pi}_i)} \right) \right)$$

The null deviance is the deviance test statistic of the model with intercept only. Residual deviance is the deviance test statistic of the model you fitted.

10.5 Diagnostic Measures for Logistic Regression

For examining adequacy of the fitted model, we will only consider the ungrouped case. Residuals can be used to diagnose adequacy. The ordinary residuals are defined as

$$e_i = Y_i - \hat{\pi}_i$$

These do not have constant variance. The deviance residuals are

$$d_i = \pm \left\{ 2 \left[Y_i \ln \left(\frac{Y_i}{n_i \hat{\pi}_i} \right) + (n_i - Y_i) \ln \left(\frac{n_i - Y_i}{n_i (1 - \hat{\pi}_i)} \right) \right] \right\}^{1/2}$$

where the sign of d_i is the same as e_i . We compute the standardized Pearson residuals as

$$r_i = \frac{Y_i - \hat{\pi}_i}{\sqrt{\hat{\pi}_i(1 - \hat{\pi}_i)}}$$

which do not have unit variance. The studentized deviance Pearson residuals are

$$sr_i = \frac{r_i}{\sqrt{1 - h_{ii}}}$$

where h_{ii} is the i th diagonal entry of the hat matrix,

$$H = V^{1/2} X (X' V X)^{-1} V^{1/2}$$

where V is the diagonal matrix with $V_{ii} = n_i \hat{\pi}_i (1 - \hat{\pi}_i)$. We can plot the deviance and the studentized Pearson residuals to check for outliers. A normal probability plot of the deviance residuals can be used to check for the fit of the model and for outliers. A plot of the deviance vs the estimated probability of success can be used to determine where the model is poorly fitted, at high or low probabilities.

For an adequate model, $E(Y_i) = \hat{\pi}_i$, and the plots of sr_i vs $\hat{\pi}_i$ and sr_i vs linear predictors $X_i' \beta$ should show a smooth horizontal Lowess line through 0. Similarly for a plot of d_i vs $\hat{\pi}_i$ and d_i vs $X_i' \beta$.

10.6 Detection of Influential Observations

In order to flag influential cases, we consider deleting one observation at a time and measuring its effects on both the χ^2 statistic and the Dev statistic. We plot these vs i , and look for spikes which indicate influential observations. Similarly, we can plot these vs $\hat{\pi}_i$.

10.6.1 Influence on the Fitted Linear Predictor

Cook's distance here measures the standardized change in the linear predictor $X' \beta$ when the i th case is deleted. Indexed plots of Cook's distance identify cases that have a large influence on the fitted predictor. Indexed plots of the leverage values h_{ii} help to identify outliers in the X space. In all cases, visual assessments are needed because there is no actual rule for flagging outlier or influential cases.

10.7 Poisson Regression

In Poisson regression, we have counting data Y which follows a Poisson distribution with mean μ , and variance μ ,

$$f(y) = \frac{e^{-\mu} \mu^y}{y!}$$

The model is then given as

$$Y_i = \mu_i + \epsilon_i, \quad i = 1, \dots, n$$

We assume there is a function g which relates the mean of the response to a linear predictor,

$$g(\mu_i) = \eta_i = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

The function g is usually called the link function. The relationship between the mean and the linear predictor is

$$\mu_i = g^{-1}(\eta_i) = g^{-1}(X_i' \beta)$$

There are several common link functions that are used with the Poisson distribution, one of them is the identity link,

$$g(\mu_i) = \mu_i = X_i' \beta$$

Another popular function is the log link,

$$g(\mu_i) = \ln(\mu_i) = X_i' \beta$$

The relationship between the mean of the response and the linear predictor is then

$$\mu_i = g^{-1}(X_i' \beta) = e^{X_i' \beta}$$

Inference is done on the Poisson model the same as in the case of logistic regression.

10.8 Generalized Linear Models

Both the logistic regression and Poisson regression models are particular examples of a more general linear model (GLM). The response is assumed to have a distribution which is a member of the exponential family of distributions,

$$f(y_i, \theta_i, \phi) = \exp \left(\frac{y_i \theta - b(\theta_i)}{a(\phi)} + h(y_i, \phi) \right)$$

Here, $\mu = E(Y) = \frac{db(\theta_i)}{d\theta_i}$, and $\text{Var}(Y) = \frac{d^2b(\theta_i)}{d\theta_i^2} a(\phi)$. The basic idea is to develop a linear model for a function of the mean. Set $\eta_i = g(\mu_i) = X_i' \beta$. The function g is the link function, for logistic regression

$$f(y_i, \theta_i, \phi) = \binom{n_i}{y_i} p^{y_i} (1-p)^{n-y_i} = \exp \left(y_i \ln \left(\frac{p}{1-p} \right) + n \ln(1-p) + \ln \binom{n_i}{y_i} \right)$$

For Poisson regression,

$$f(y_i, \theta_i, \phi) = \frac{e^{-\lambda} \lambda^{y_i}}{y_i!} = \exp(y_i \ln \lambda - \lambda - \ln(y_i!))$$